
Lecture Notes for Machine Learning in Python

Professor Eric Larson
Recurrent Neural Networks

Lecture Agenda

- Logistics
 - CNNs due yesterday
 - RNNs due next week (Friday)
- Recurrent Networks (two lecture agenda)
 - Overview
 - Problem Types
 - Embeddings
 - Types of RNNs
 - Demo A
 - CNNs and RNNs
 - Demo B
 - course retrospective

History of Recurrent Neural Networks

LIFE SCORECARD

TIMES WHEN I THOUGHT...

"I'M NOT REALLY HAPPY HERE, BUT
MAYBE THIS IS THE BEST I CAN EXPECT
AND I'LL REGRET GIVING IT UP."

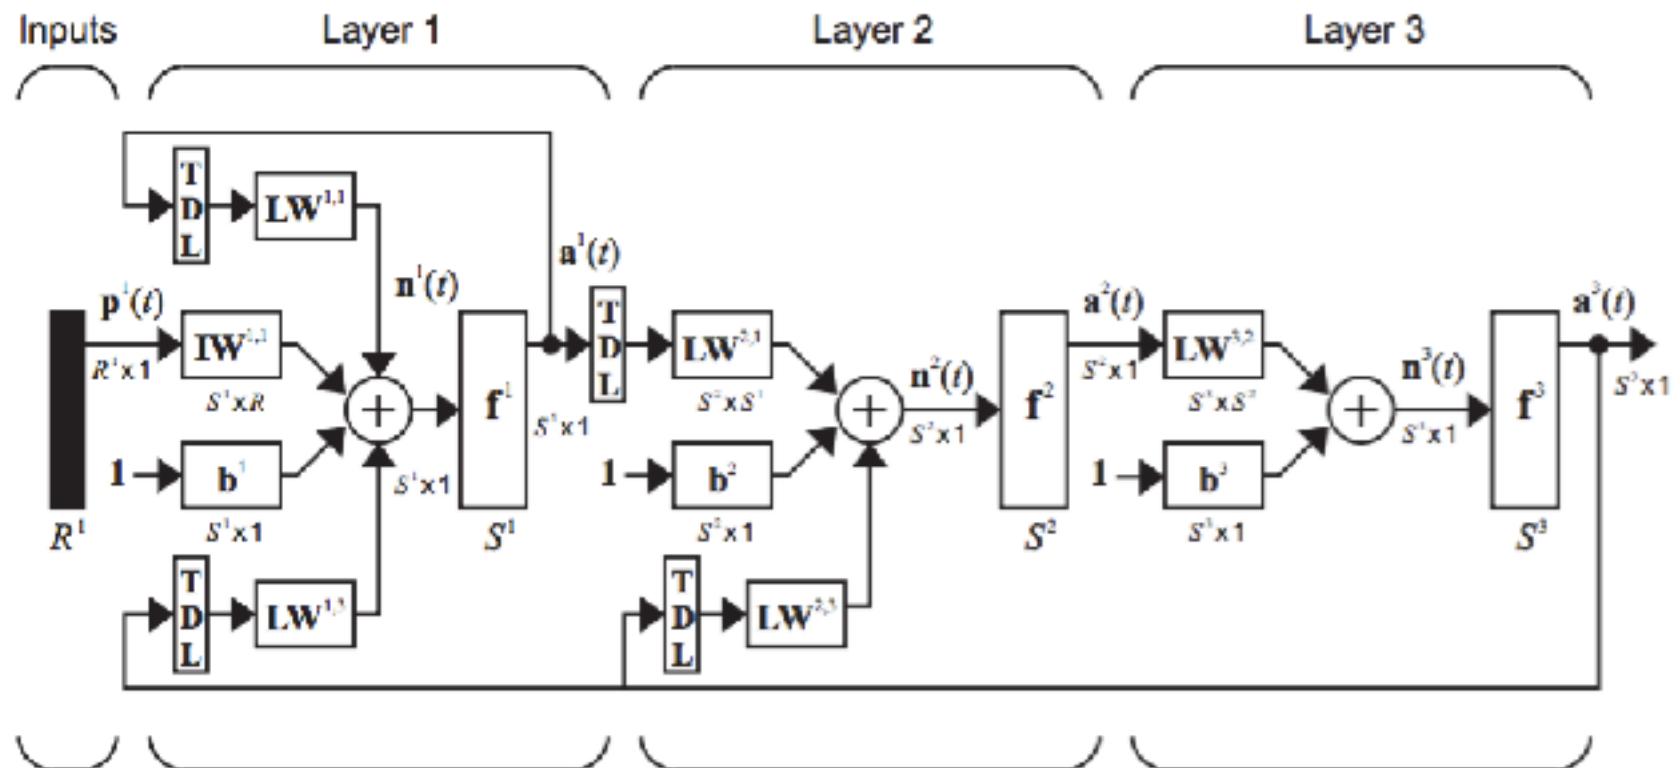
...IT TURNED OUT I...

SHOULD HAVE STAYED	SHOULD HAVE LEFT SOONER
--------------------	-------------------------

II	HTH HTH III
----	-------------

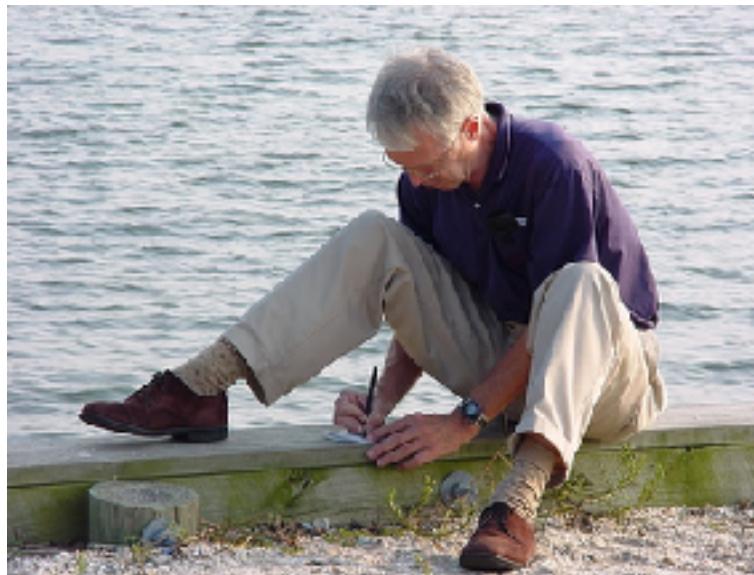
History of Recurrent Networks

- Dynamic Networks
 - can use current and previous inputs, in time
 - still popular, but ultimately extremely hard to train
 - **been around for decades**

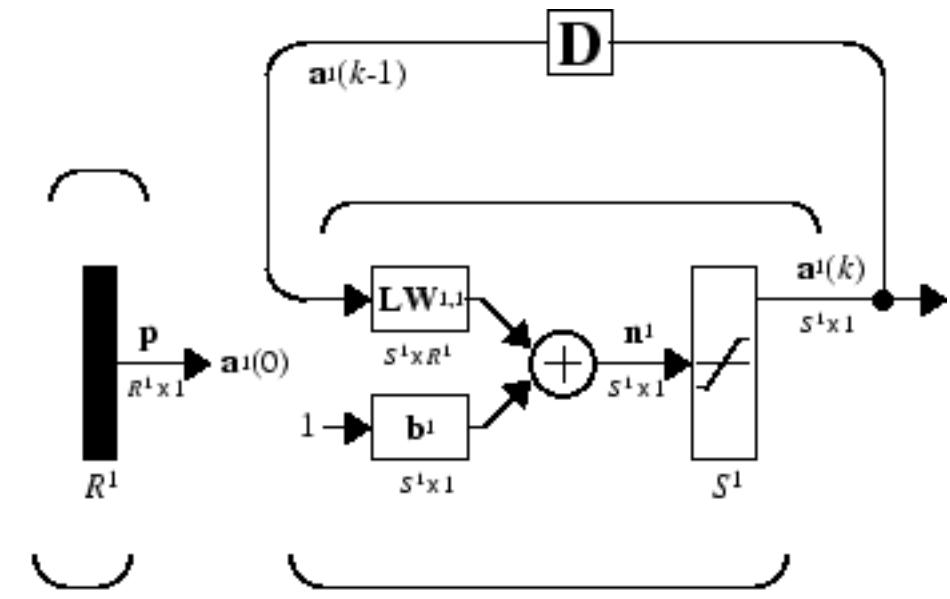


History of Recurrent Networks

- Hopfield Network, 1982



John Hopfield, Princeton



Initial conditions

Symmetric saturated linear layer

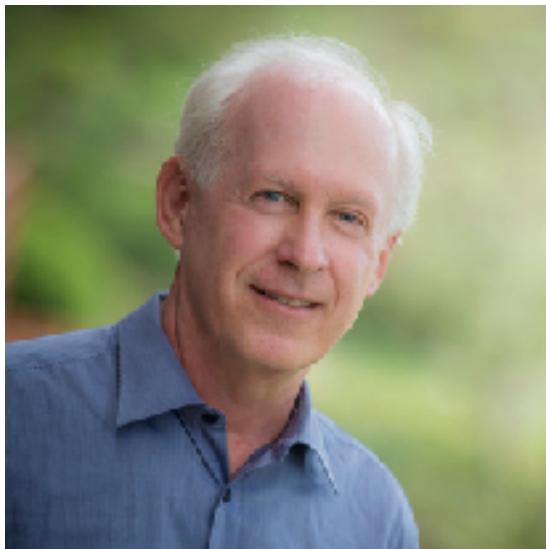
$\mathbf{a}^1(0) = \mathbf{p}$ and then for $k = 1, 2, \dots$

$\mathbf{a}^1(k) = \text{satlims}(\mathbf{LW}^{1,1}\mathbf{a}^1(k-1)) + \mathbf{b}^1$

Neural Network Design, Hagan, Demuth, Beale, and De Jesus

History of Recurrent Networks

- Elman/Jordan Networks, ~1988



Jeffrey Elman, UCSD



Michael Jordan, Berkeley

Elman network^[10]

$$h_t = \sigma_h(W_h x_t + U_h h_{t-1} + b_h)$$

$$y_t = \sigma_y(W_y h_t + b_y)$$

Jordan network^[11]

$$h_t = \sigma_h(W_h x_t + U_h y_{t-1} + b_h)$$

$$y_t = \sigma_y(W_y h_t + b_y)$$

Variables and functions

- x_t : input vector
- h_t : hidden layer vector
- y_t : output vector
- W , U and b : parameter matrices and vector
- σ_h and σ_y : Activation functions

Wikipedia

History of Recurrent Networks

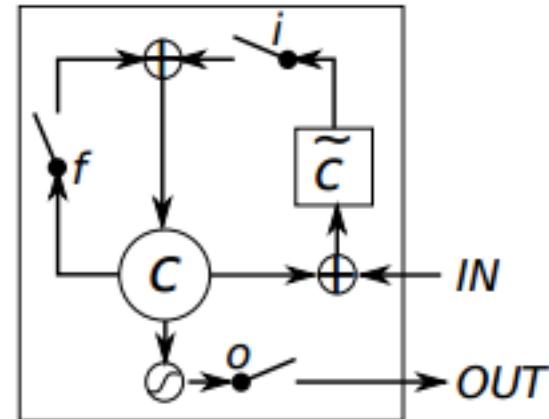
- Long Short Term Memory, ~1997



Sepp Hochreiter, Many Universities



Jürgen Schmidhuber, Switzerland



More on these later

History of Recurrent Networks

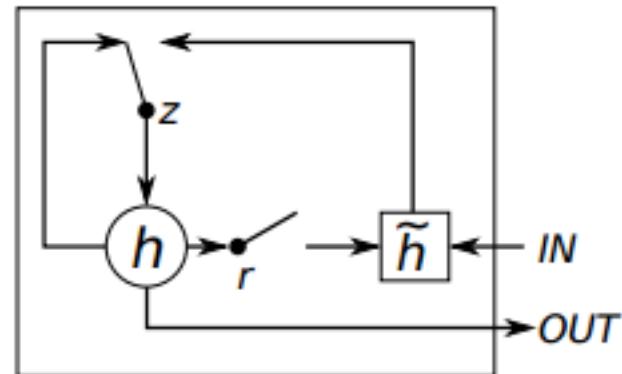
- Gated Recurrent Units, ~2014



Yoshua Bengio



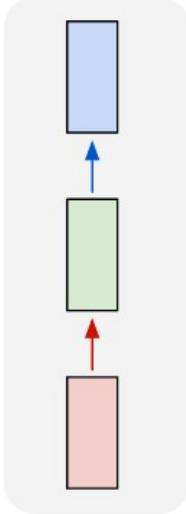
Kyunghyun Cho, Professor at NYU



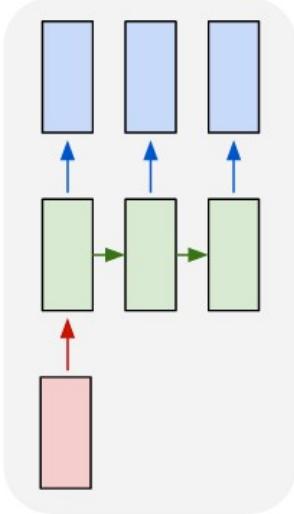
More on these later

Recurrent Networks: Problem Types

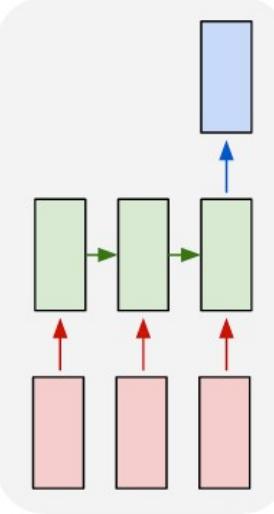
one to one



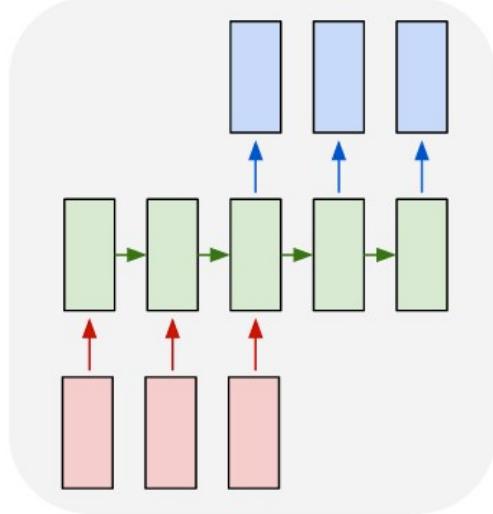
one to many



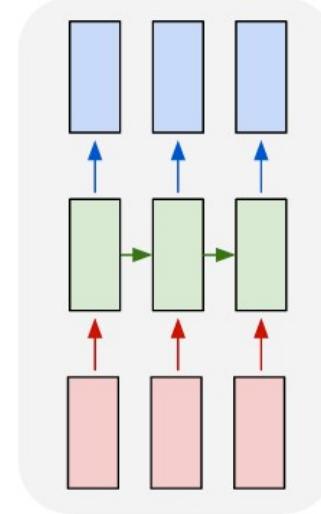
many to one



many to many

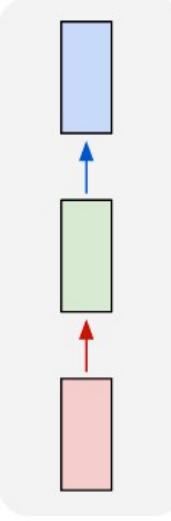


many to many

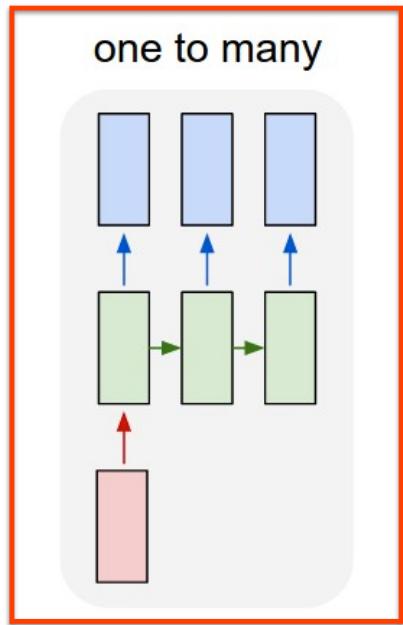


Recurrent Networks: Problem Types

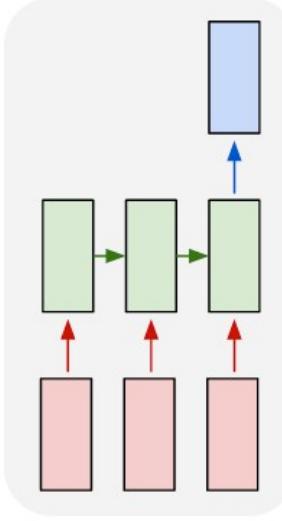
one to one



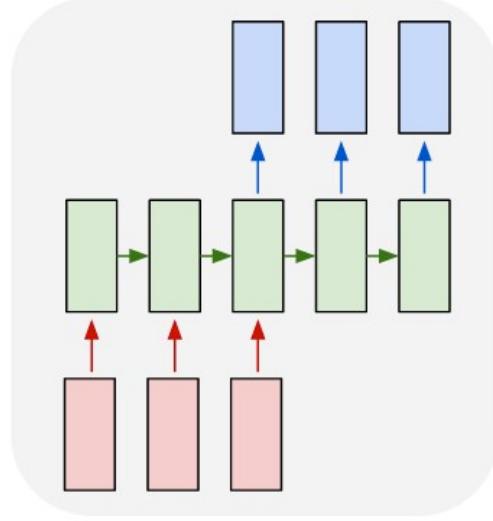
one to many



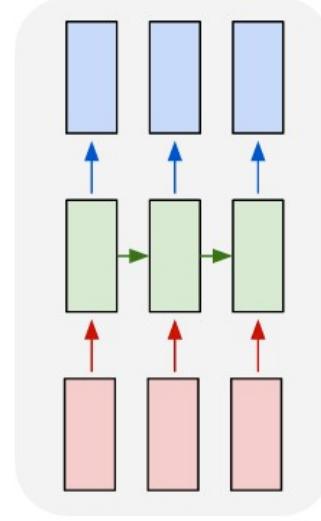
many to one



many to many



many to many



A person riding a motorcycle on a dirt road.



Two dogs play in the grass.



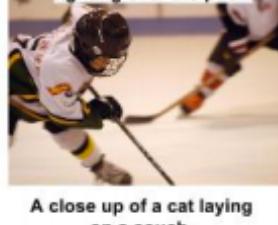
A skateboarder does a trick on a ramp.



A group of young people playing a game of frisbee.



Two hockey players are fighting over the puck.



A little girl in a pink hat is blowing bubbles.



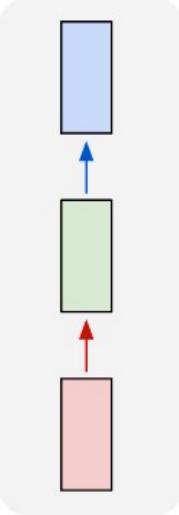
A herd of elephants walking across a dry grass field.

A close up of a cat laying on a couch.

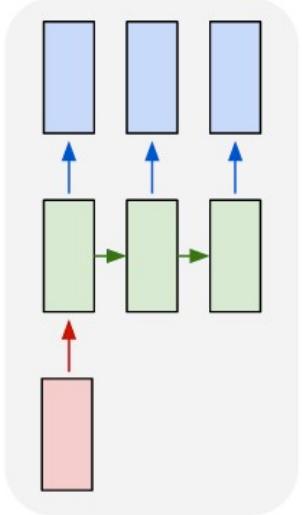
A red motorcycle parked on the side of the road.

Recurrent Networks: Problem Types

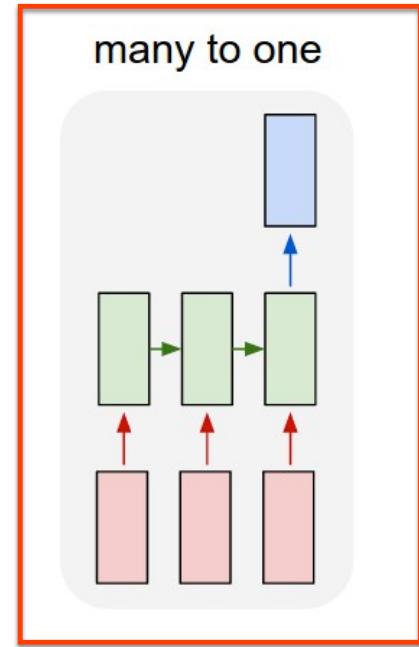
one to one



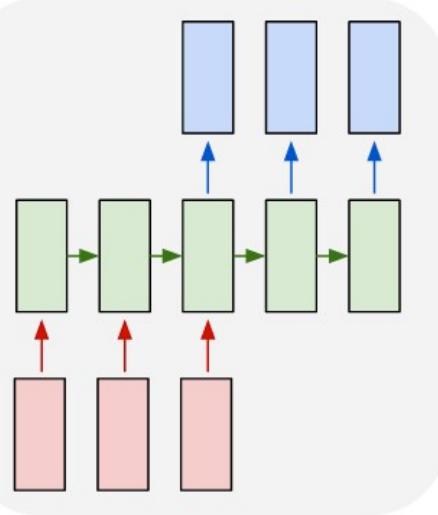
one to many



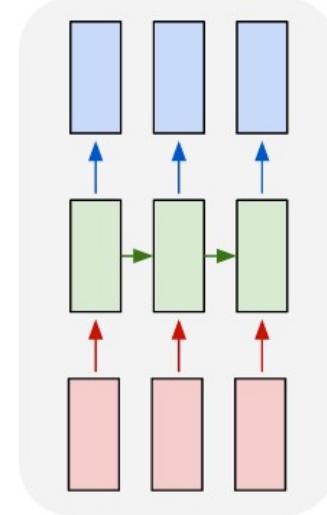
many to one



many to many



many to many



The movie is great.



The movie stars Mr. X



The movie is horrible.



Recurrent Networks: Ontology Classification

Eva Ingolf is a well known Icelandic violinist particularly recognized for her authoritative performances of solo works by J. S. Bach. She comes from a leading musical family and her father Ingólfur Guðbrandsson premiered many of the great choral works in Iceland and six of her sisters and brothers are professional musicians who have made an important contribution to the high quality of the musical life in the country. Eva Ingolf currently lives in New York City with her husband Kristinn Sv.

Artist

Shaun Norris (born 14 May 1982) is a South African professional golfer. Norris plays on the Sunshine Tour where he has won twice. He won the inaugural Africa Open in 2008 and the Nashua Masters in 2011. He also began playing on the European Tour in 2011 after graduating from qualifying school.

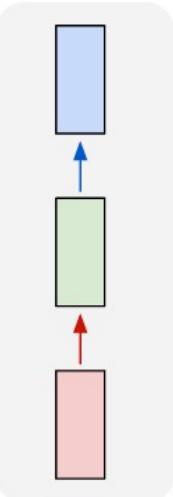
Athlete

Palace Software was a British video game publisher and developer during the 1980s based in London England. It was notable for the Barbarian and Cauldron series of games for 8-bit home computer platforms in particular the ZX Spectrum Amstrad CPC and Commodore 64.

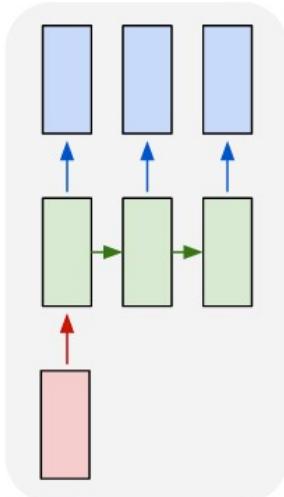
Company

Recurrent Networks: Problem Types

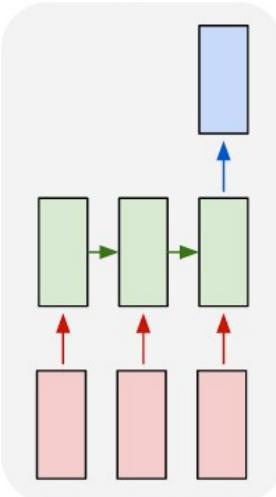
one to one



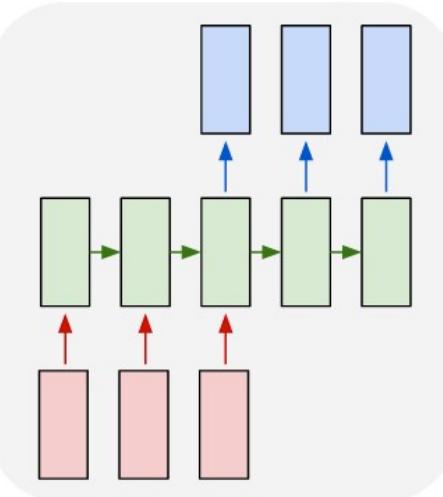
one to many



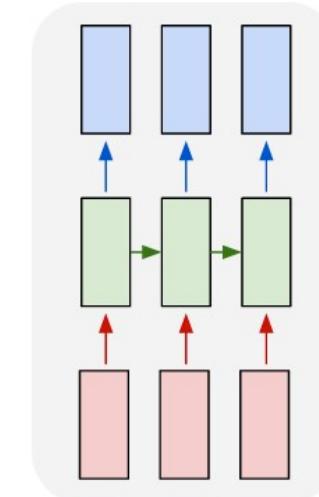
many to one



many to many



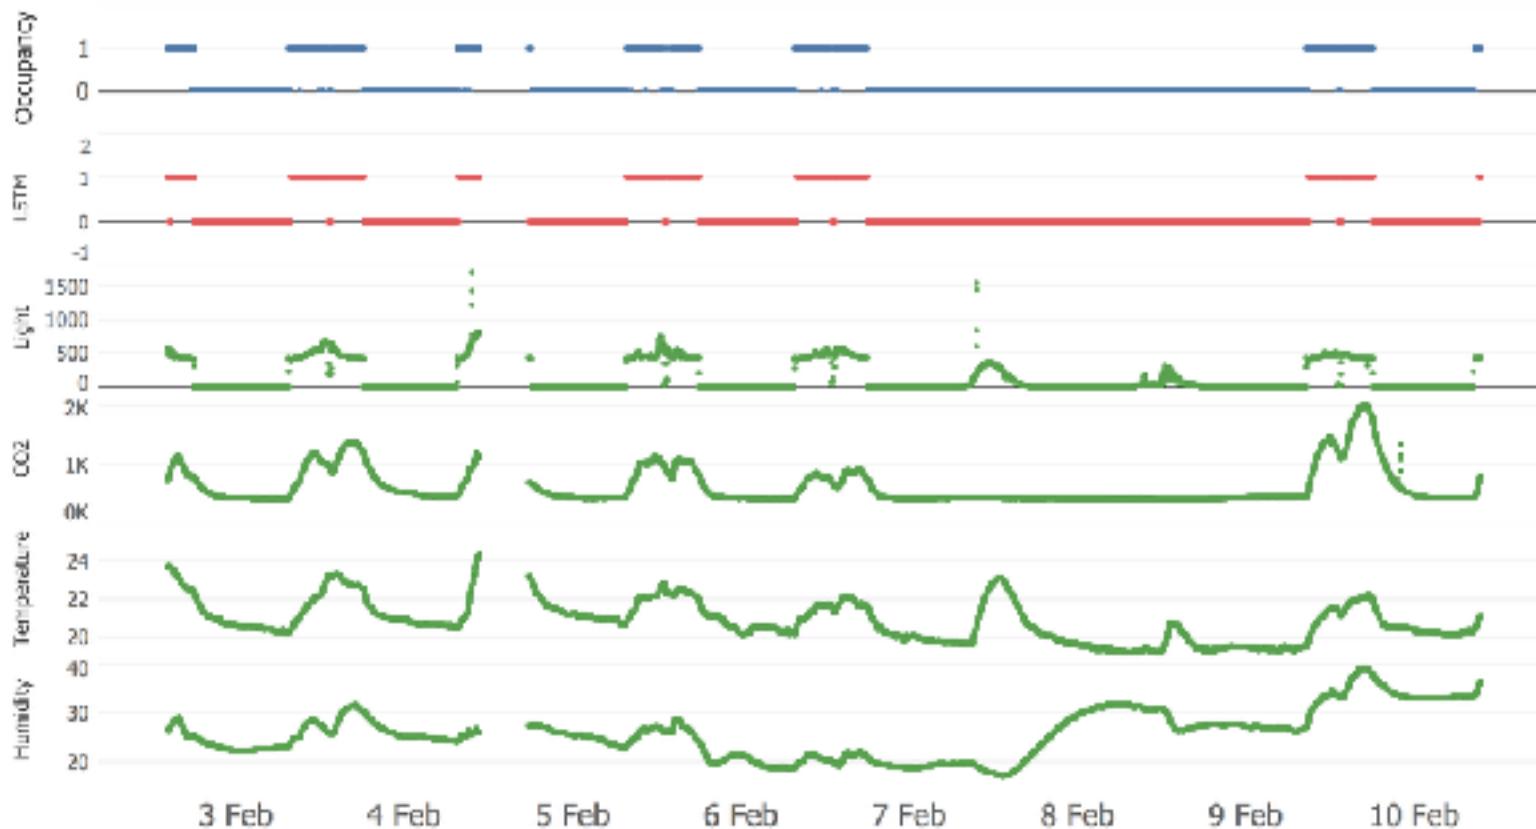
many to many



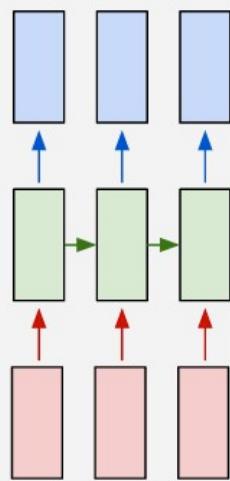
Das Wirtschaftswachstum hat sich in den letzten Jahren verlangsamt .
Economic growth has slowed down in recent years .

La croissance économique s' est ralentie ces dernières années .

Recurrent Networks: Problem Types

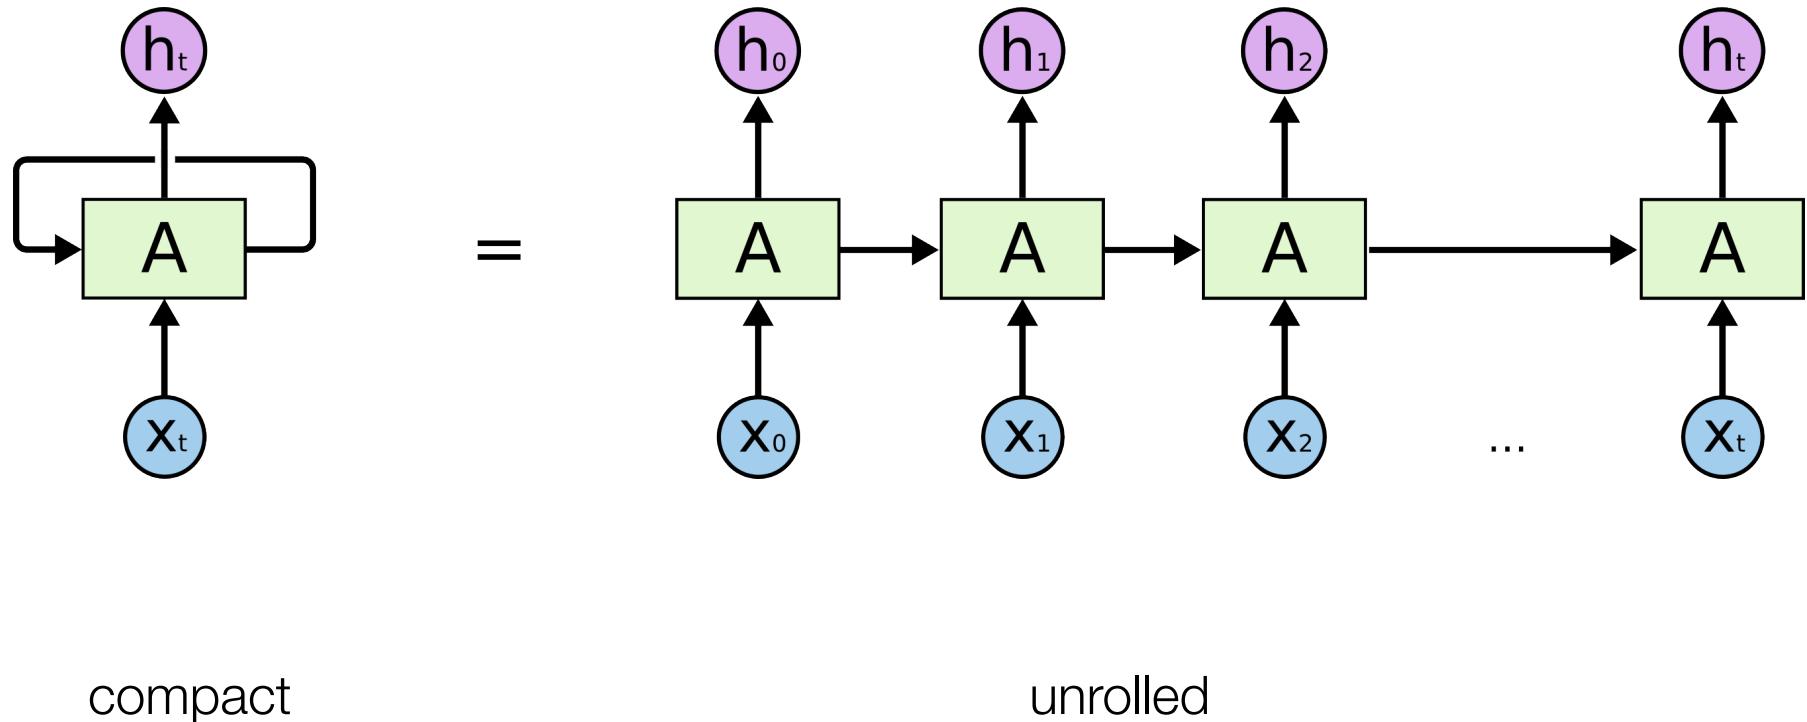


many to many

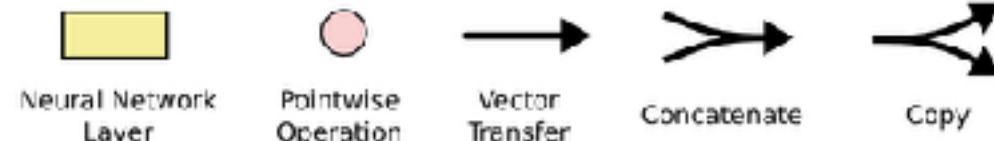


sequence to sequence

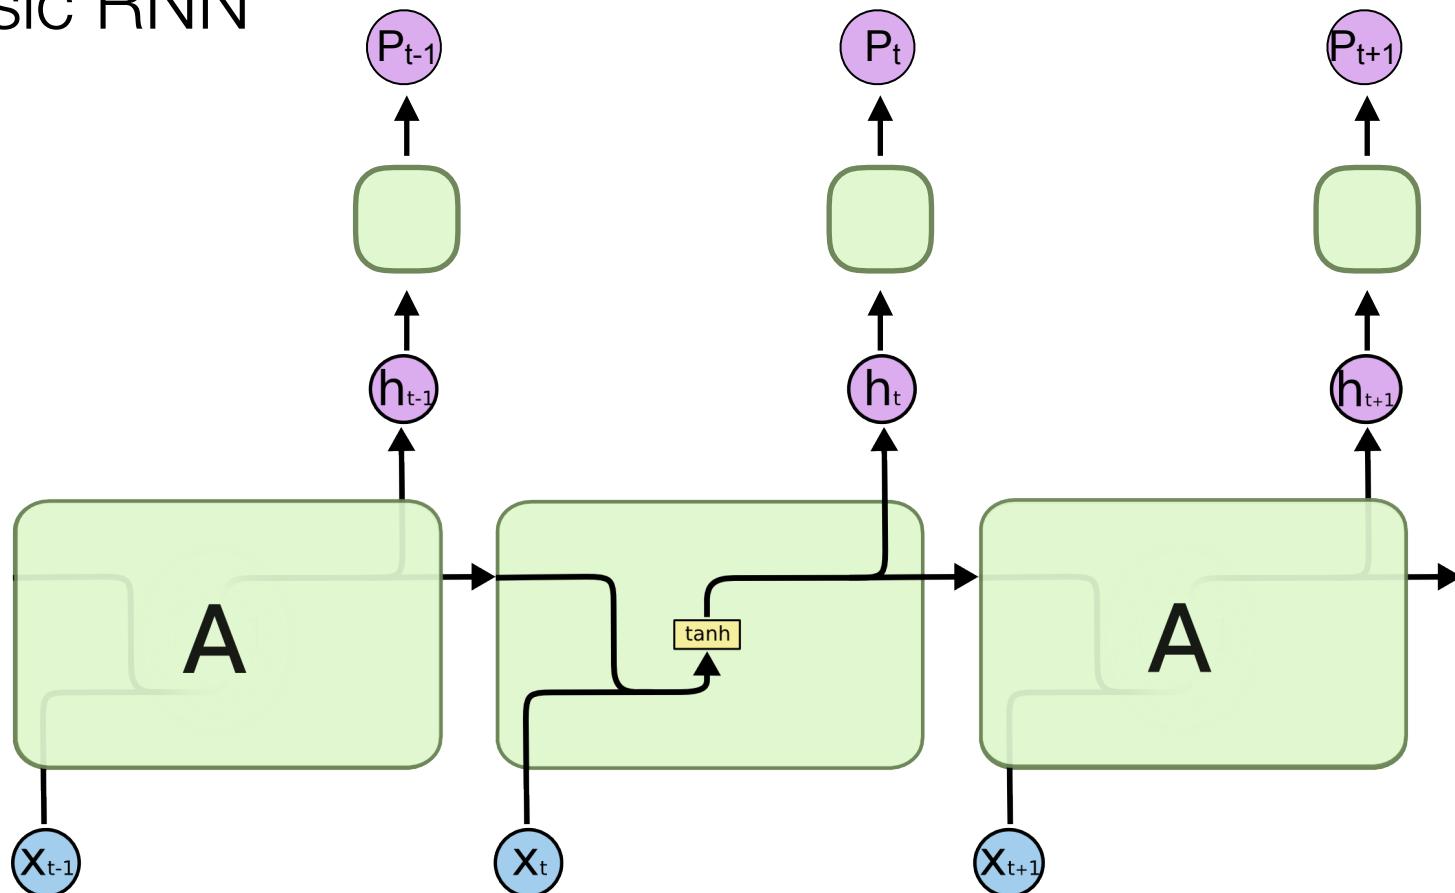
Recurrent Networks: Main Idea



Recurrent Networks



- basic RNN

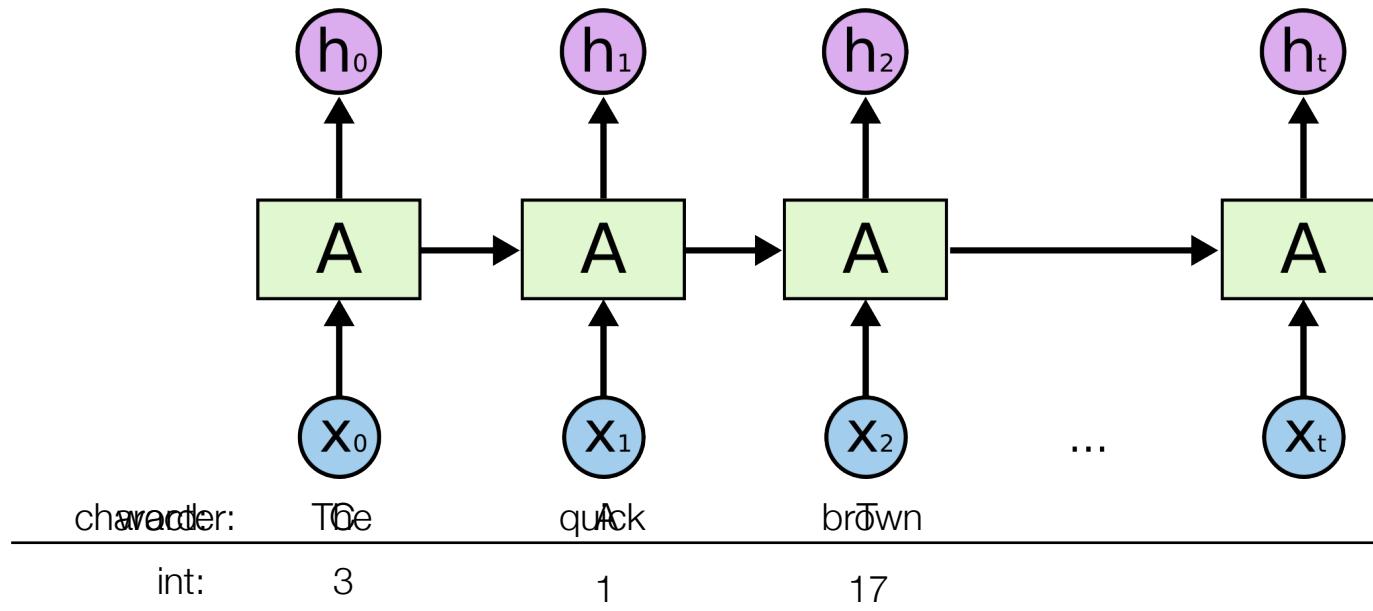


$$h_t = \tanh(W_A (X_t \oplus h_{t-1}) + b_A)$$

$$P_t = \text{softmax}(W_P h_t + b_P)$$

Recurrent Networks: Representation

- python:



...

0

0

...

0

0

...

0

0

...

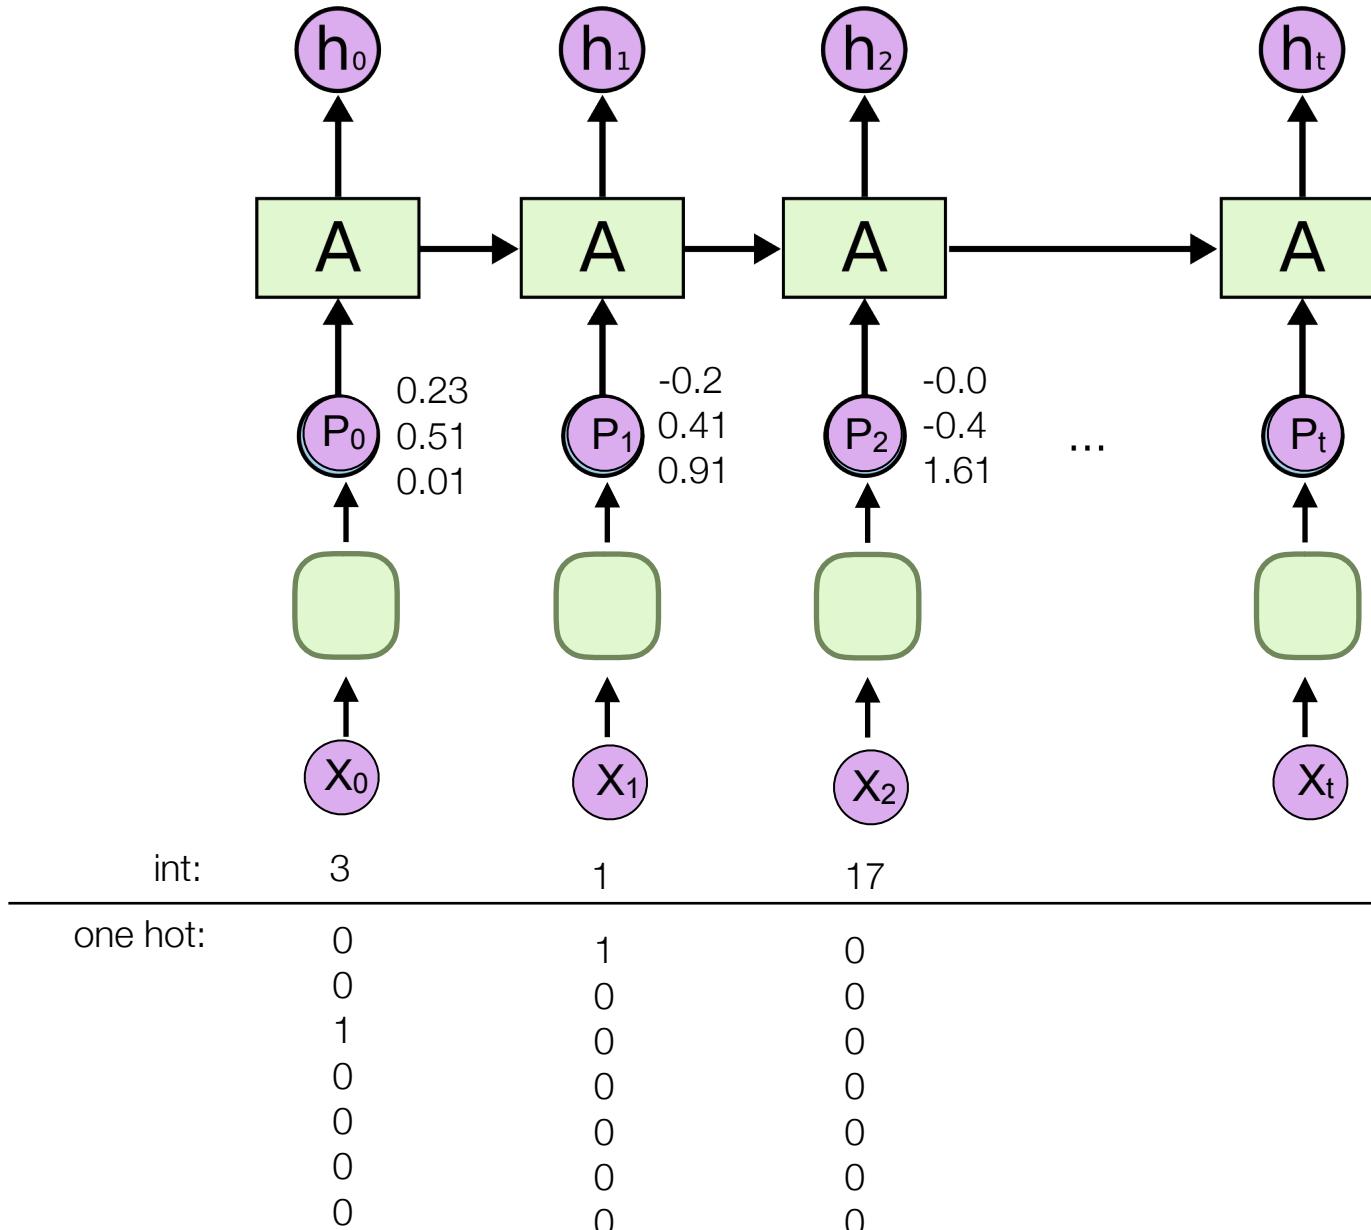
1

0

<http://colah.github.io/posts/2015-08-Understanding-LSTMs/>

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>

Word Embeddings



Word Embeddings • Many are pre-trained for you!!

GloVe

Global Vectors for Word Representation

Highlights

1. Nearest neighbors

The Euclidean distance (or cosine similarity) between two word vectors provides an effective method for measuring the linguistic or semantic similarity of the corresponding words. Sometimes, the nearest neighbors according to this metric reveal rare but relevant words that lie outside an average human's vocabulary. For example, here are the closest words to the target word *frog*:

- 0. *frog*
- 1. *frogs*
- 2. *toad*
- 3. *litoria*
- 4. *leptodactylidae*
- 5. *rana*
- 6. *lizard*
- 7. *eleutherodactylus*



3. *litoria*



4. *leptodactylidae*

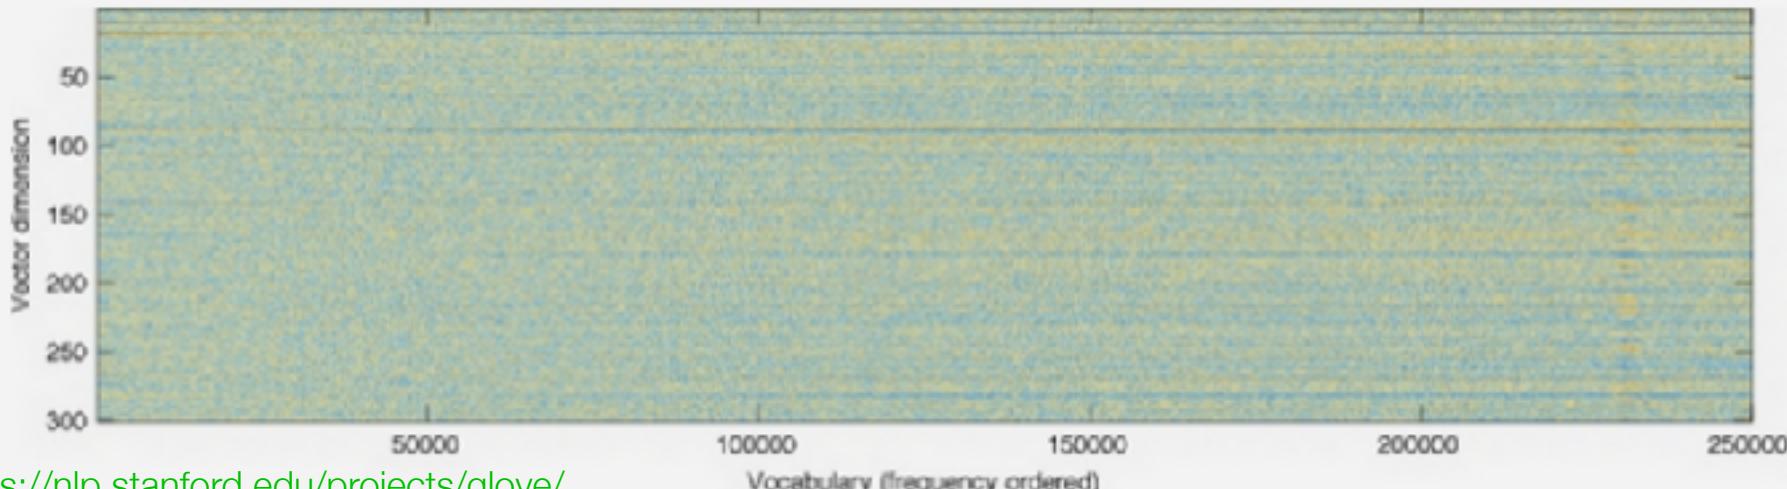


5. *rana*



7. *eleutherodactylus*

GloVe produces word vectors with a marked banded structure that is evident upon visualization:

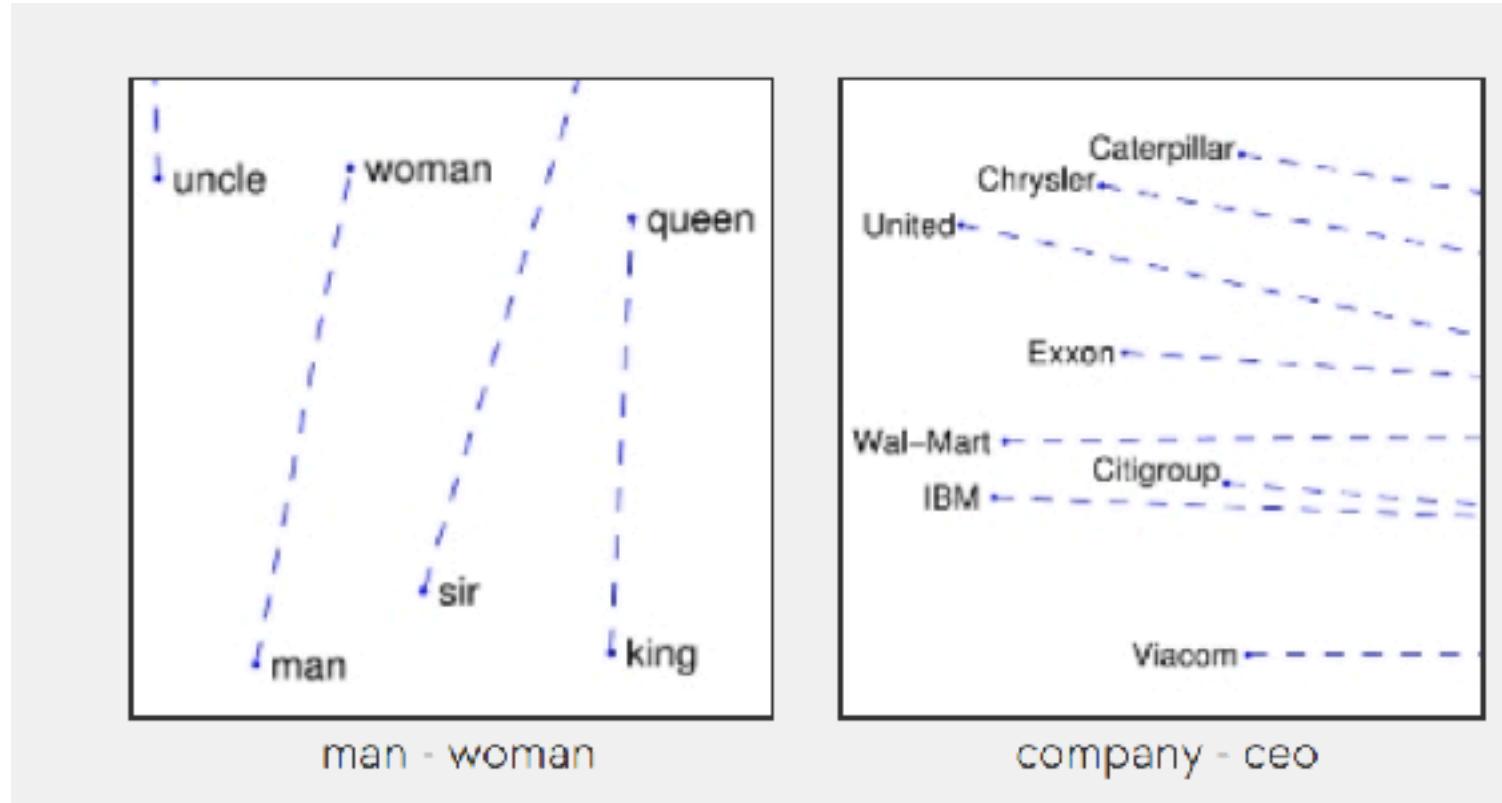


Word Embeddings

• Many are pre-trained for you!!

GloVe

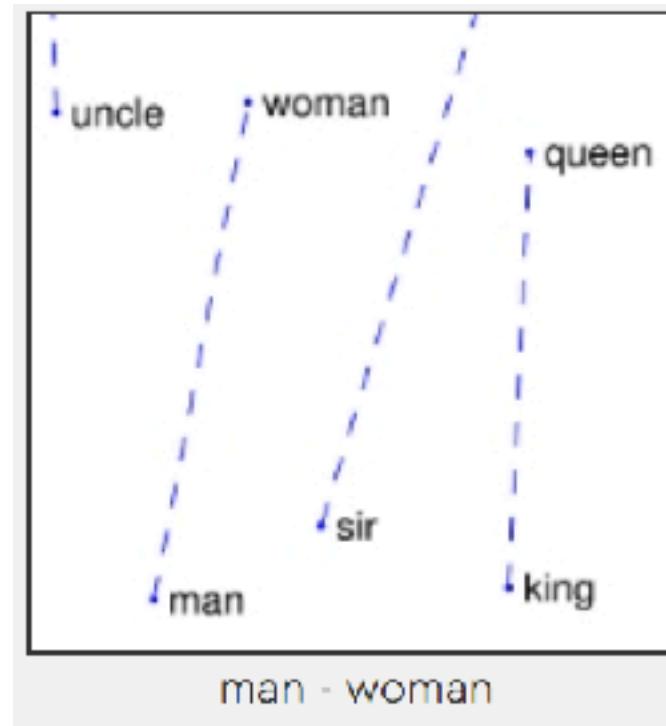
Global Vectors for Word Representation



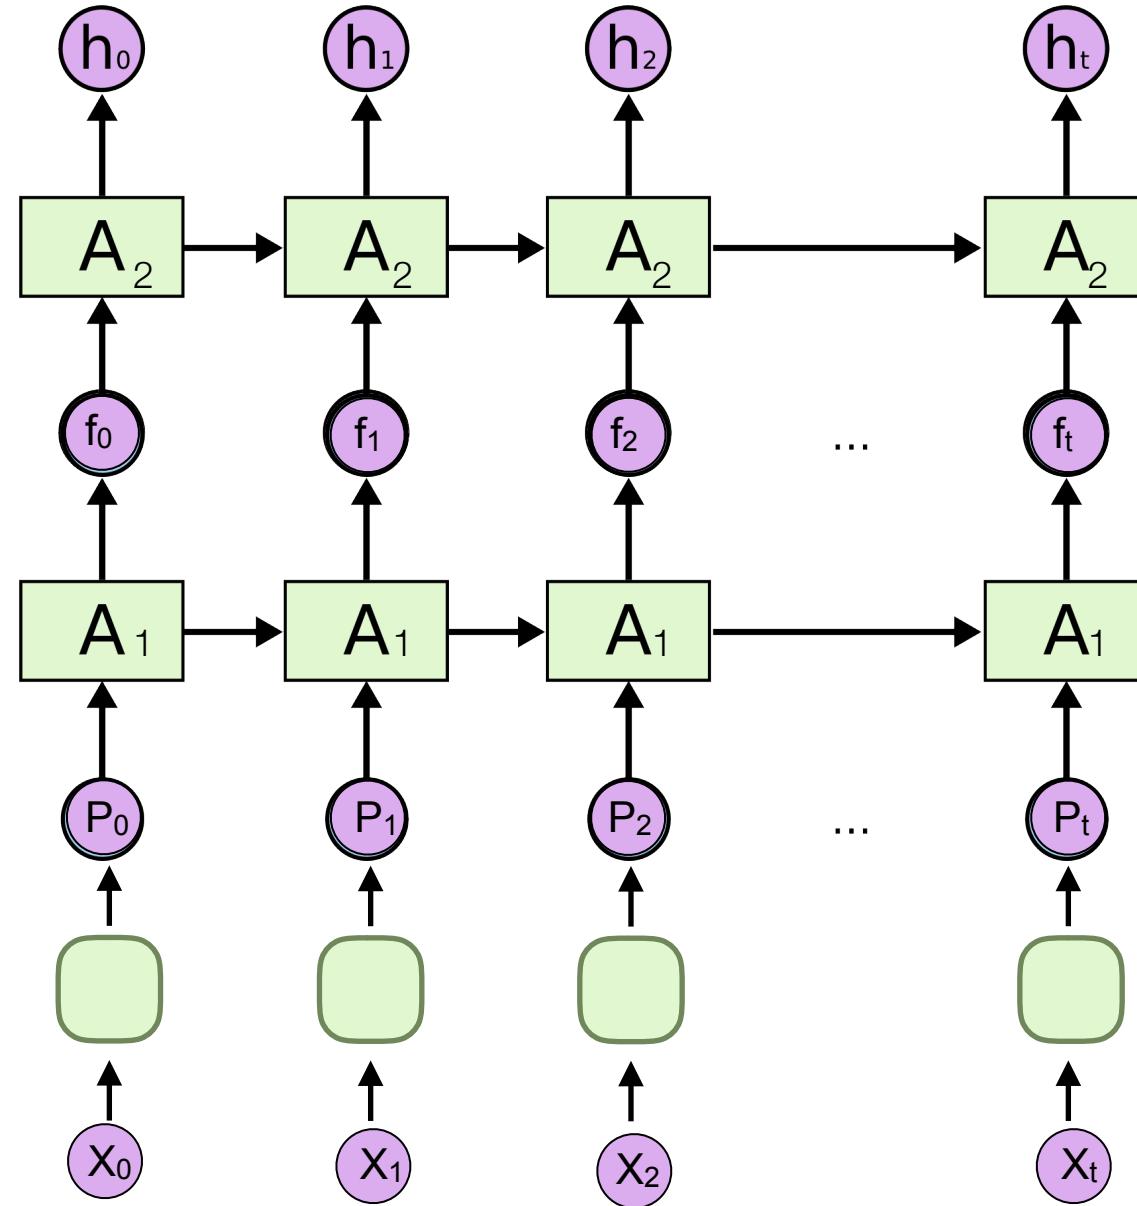
each axis might encode a different type of relationship

Self Test

- Each axis on the embedding plot below corresponds to:
 - a weight inside the embedding layer
 - the average of weights in the embedding layer
 - the average of the one hot encoding for a word
 - an output of the embedding layer

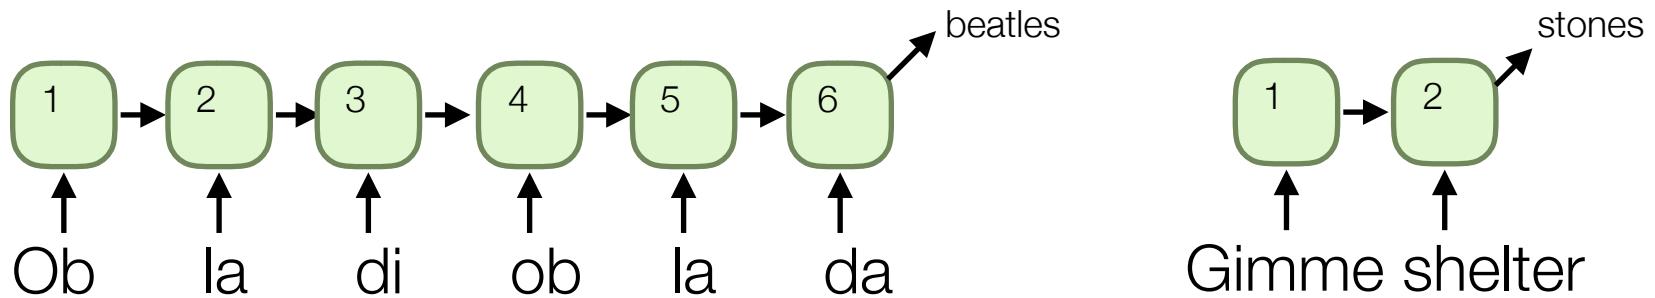


Sequence Stacking



Practical Logistics: Sequence Length

- option A: dynamic length sequences



- option B: padding/clipping

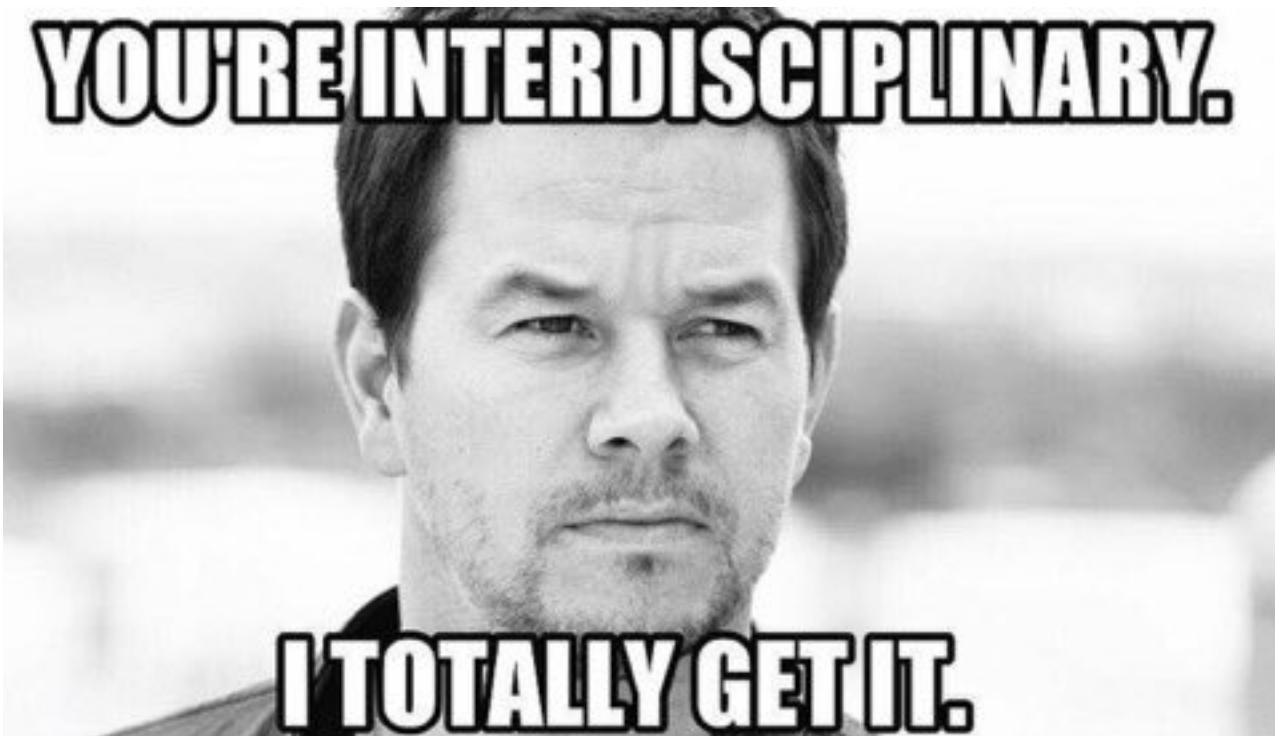


- main difference: **speed**

Self Test

- The main reason dynamic length is slow is because:
 - A. the computation graph must be updated
 - B. weights must be tied together for each recurrent node
 - C. the weights must be multiplied until the output converges
 - D. the unrolling operation takes some time

For Fun: Recurrent Generation



YOU'RE INTERDISCIPLINARY.

I TOTALLY GET IT.

Generating More Outputs

- Highly sophisticated steps:
 - train an RNN to generate the **next** word/character from the **current** word/character
 - train on a corpus of text
 - Shakespeare
 - Movie Scripts
 - Whatever!
 - seed with random word,
feed output words as input to next node
 - rinse, repeat

Recurrent Networks: Character Data

First Citizen:

Before we proceed any further, hear me speak.

All:

Speak, speak.

First Citizen:

You are all resolved rather to die than to famish?

All:

Resolved. resolved.

First Citizen:

First, you know Caius Marcius is chief enemy to the people.

All:

We know't, we know't.

First Citizen:

Let us kill him, and we'll have corn at our own price.

Is't a verdict?

All:

No more talking on't; let it be done: away, away!

Second Citizen:

One word, good citizens.

ervices he has done for his country?

First Citizen:

We are accounted poor citize
What authority surfeits on w
would yield us but the super
wholesome, we might guess th
but they think we are too de
afflicts us, the object of o
inventory to particularise t
sufferance is a gain to them
our pikes, ere we become rak
speak this in hunger for bre

Second Citizen:

Would you proceed especially

All:

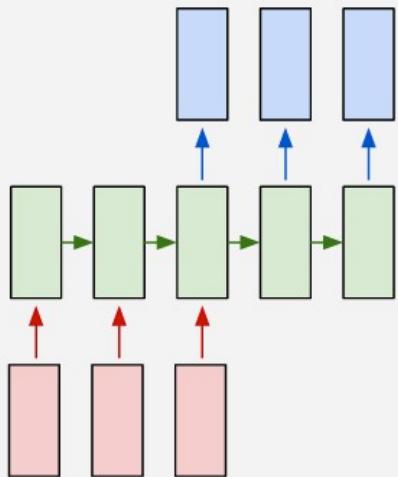
Against him first: he's a ve

Second Citizen:

Consider you what s

Recurrent Networks: Character Generation

many to many



ATOOOS

UIEAOUYOUZZZZUZAAAYAYf n fsflflrurctuateot t ta's a wtut
Whith then, a do makes and them and to sees,
I wark on this ance may string take thou honon
To sorriccorn of the bairer, whither, all
I'd see if yiust the would a peid.

LARYNGLe:

To would she troust they fould.

PENMES:

Thou she so the havin to my shald woust of
As tale we they all my forder have
As to say heant thy wansing thag and
Whis it thee shath his breact, I be and might, she
Tirs you desarvishensed and see thee: shall,
What he hath with that is all time,
And sen the have would be sectiens, way thee,
They are there to man shall with me to the mon,
And mere fear would be the balte, as time an at
And the say oun touth, thy way womers thee.

Recurrent Networks: Character Generation

more data, star wars + star trek + tarantino + the matrix

DENT ' SUEENCK

Bartholomew of the TIE FIGHTERS are stunned. There is a crowd and armored
switcheroos.

PICARD

(continuing)

Couns two dim is tired. In order to the sentence...

The sub bottle appears on the screen into a small shuttle shift of the
ceiling. The DAMBA FETT splash fires and matches them into the top, transmit to stable high above
upon their statels,
falling from an alien shaft.

ANAKIN and OBI-WAN stand next to OBI-WAN down the control plate of smoke at the TIE fighter. They
stare at the centre of the station loose into a comlink cover -- comes up to the General, the
GENERAL HUNTAN AND FINNFURMBARD from the PICADOR to a beautiful Podracisly.

ENGINEER

Naboo from an army seventy medical
security team area re-weilergular.

EXT.

THE MULTIVERSE —

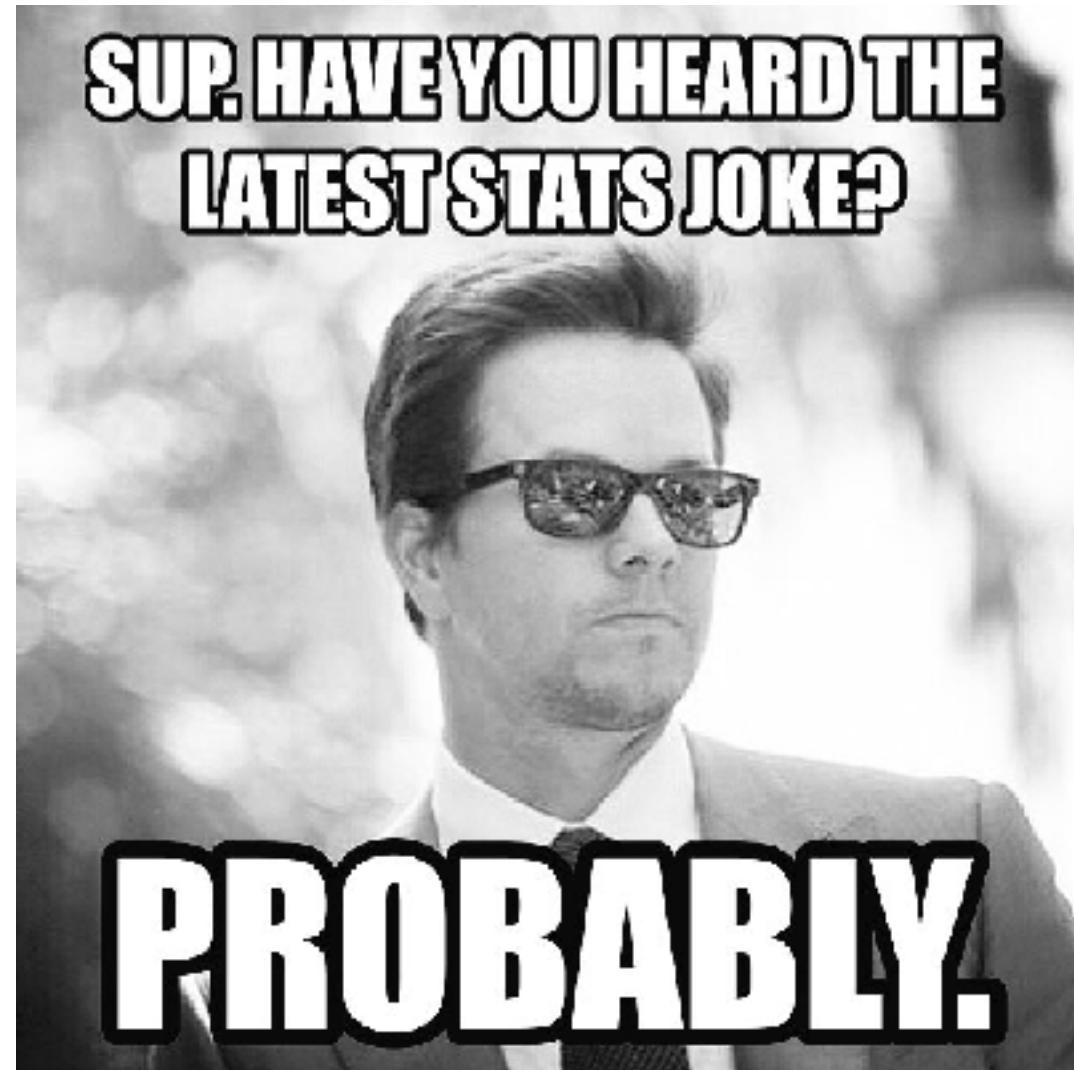
Movie written by algorithm turns out to be hilarious and intense

For *Sunspring*'s exclusive debut on Ars, we talked to the filmmakers about collaborating with an AI.

ANNALEE NEWITZ - 6/9/2016, 5:30 AM

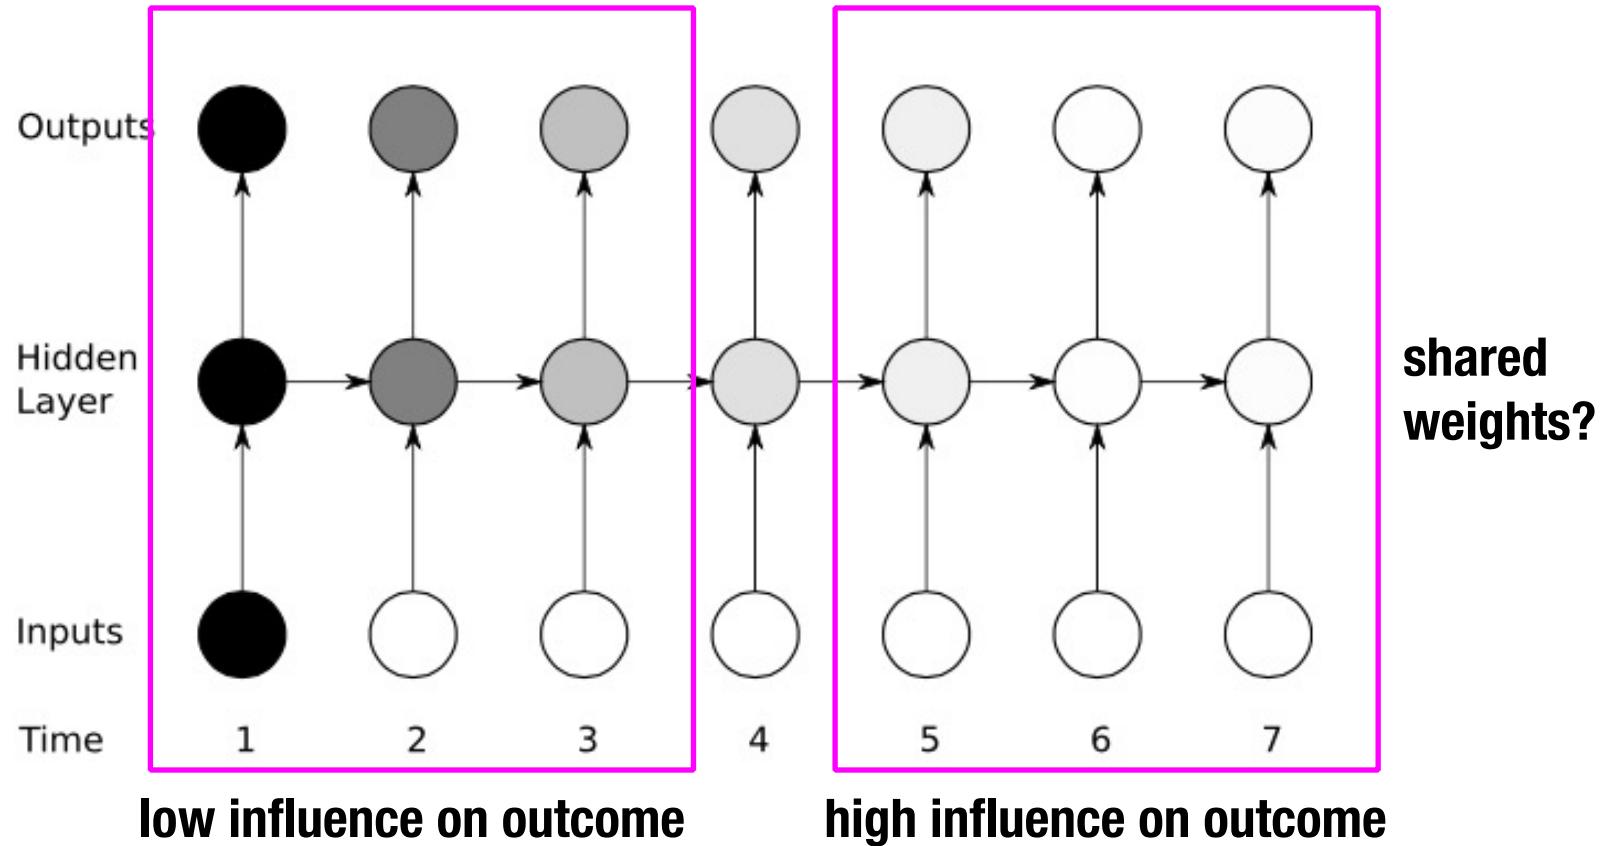


Back to Reality: Types of RNNs



Recurrent Networks, the Age Old Problem

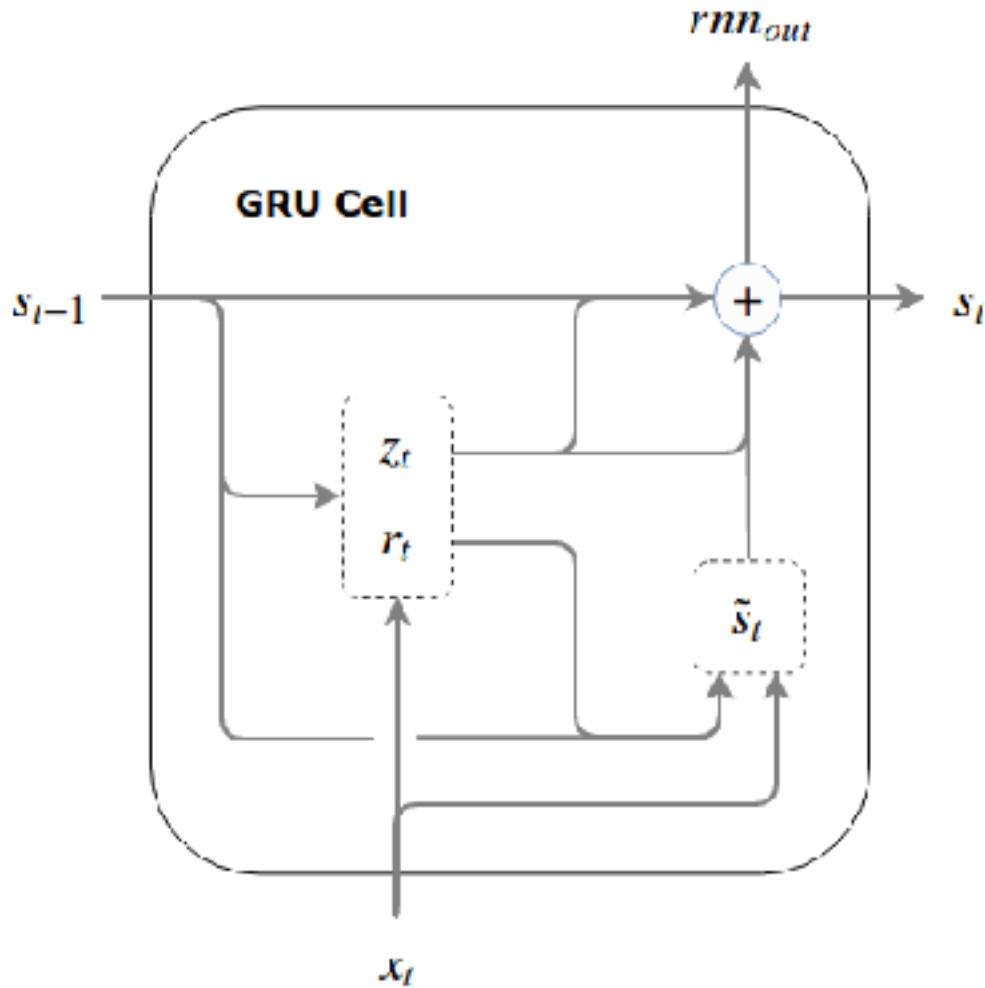
- vanishing gradients: is this the same?



$$\frac{\partial E_t}{\partial S_{t-k}} = \frac{\partial E_t}{\partial S_t} \frac{\partial S_t}{\partial S_{t-k}} = \frac{\partial E_t}{\partial S_t} \left(\frac{\partial S_t}{\partial S_{t-1}} \frac{\partial S_{t-1}}{\partial S_{t-2}} \dots \frac{\partial S_{t-k+1}}{\partial S_{t-k}} \right) = \frac{\partial E_t}{\partial S_t} \prod_{i=1}^k \frac{\partial S_{t-i+1}}{\partial S_{t-i}}$$

Recurrent Networks: GRUs

- gated recurrent units



Selectivity controls, gates (**0 or 1**)

$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$

$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

past state

current input

selectively remember

with influence

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + U x_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

remember only past

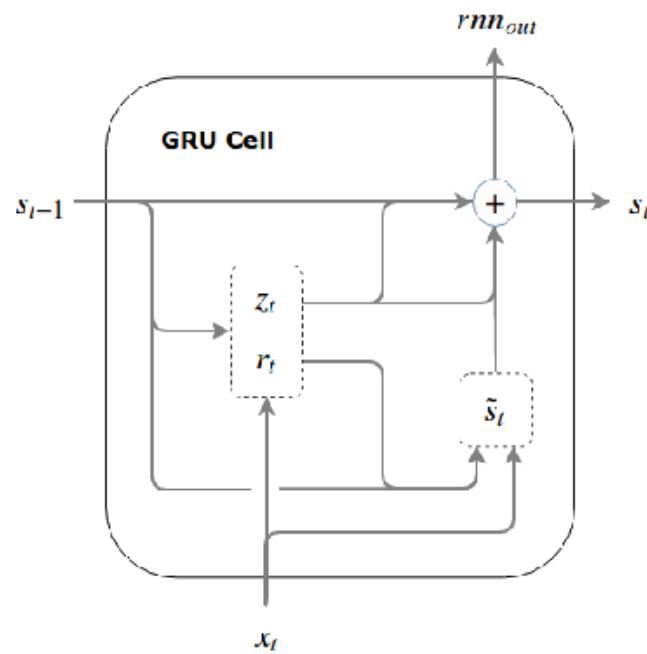
or remember with input

σ = hard limit

\odot = elem. multiplication

Self Test

- What element of the GRU helps with vanishing and exploding gradients?
- A. derivative of ϕ
- B. derivative of σ
- C. σ
- D. ϕ



$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$

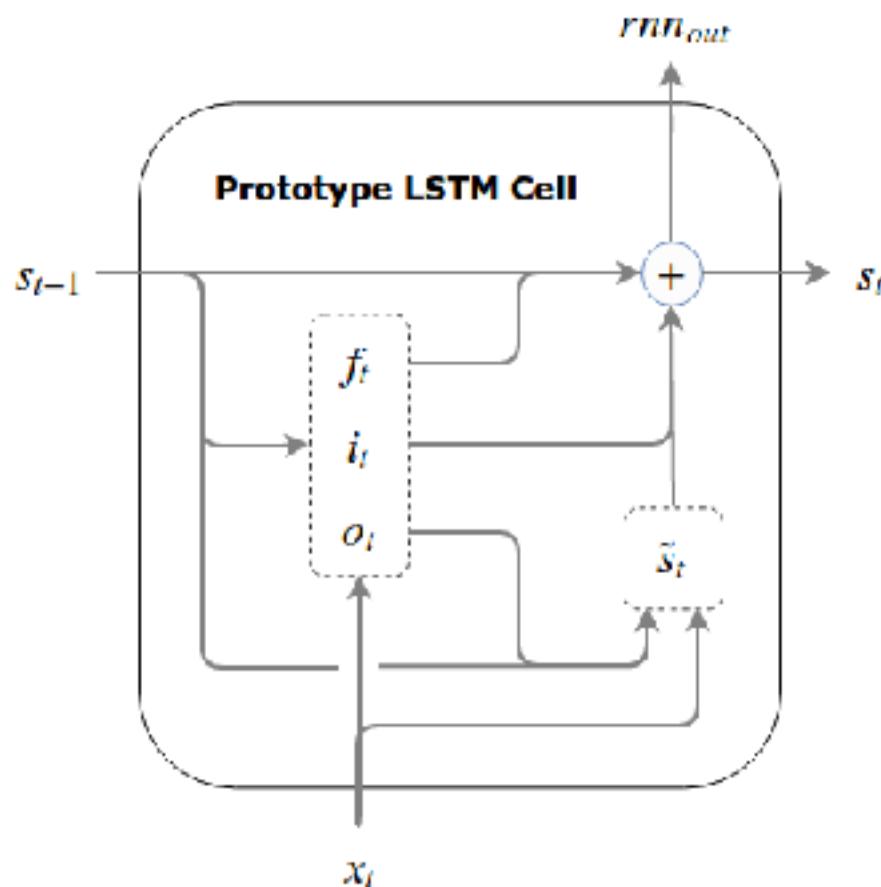
$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + Ux_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

Recurrent Networks: Gen 1 LSTM

- LSTM prototype



Selectivity controls (**gates, 0 or 1**)

$$o_t = \sigma(W_o s_{t-1} + U_o x_t + b_o)$$

$$i_t = \sigma(W_i s_{t-1} + U_i x_t + b_i)$$

$$f_t = \sigma(W_f s_{t-1} + U_f x_t + b_f)$$

selectively remember past with influence

$$\tilde{s}_t = \phi(W(o_t \odot s_{t-1}) + Ux_t + b)$$

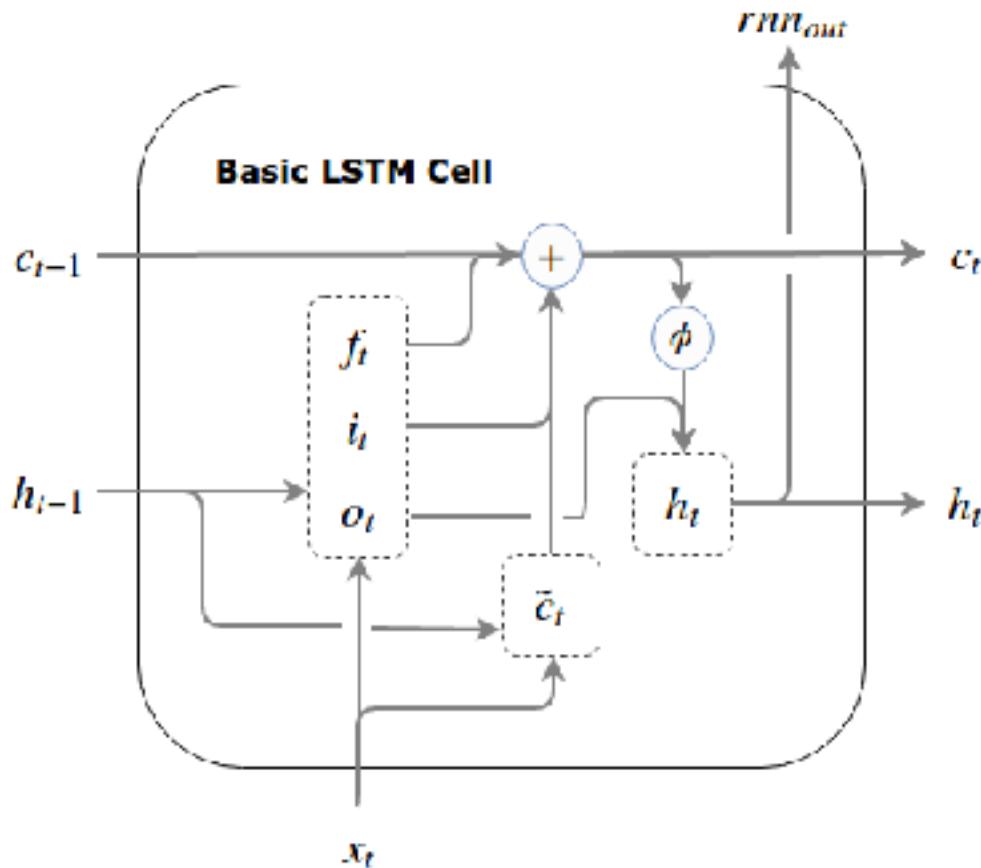
$$s_t = [f_t \odot s_{t-1}] + [i_t \odot \tilde{s}_t]$$

selectively remember past

with past weighted influence

Recurrent Networks: Gen 2 LSTM

- LSTM in TensorFlow



Selectivity controls

$$i_t = \sigma(W_i h_{t-1} + U_i x_t + b_i)$$

$$o_t = \sigma(W_o h_{t-1} + U_o x_t + b_o)$$

$$f_t = \sigma(W_f h_{t-1} + U_f x_t + b_f)$$

explicit remembering state

$$\tilde{c}_t = \phi(W h_{t-1} + U x_t + b)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

remember
previous state

update with
new state

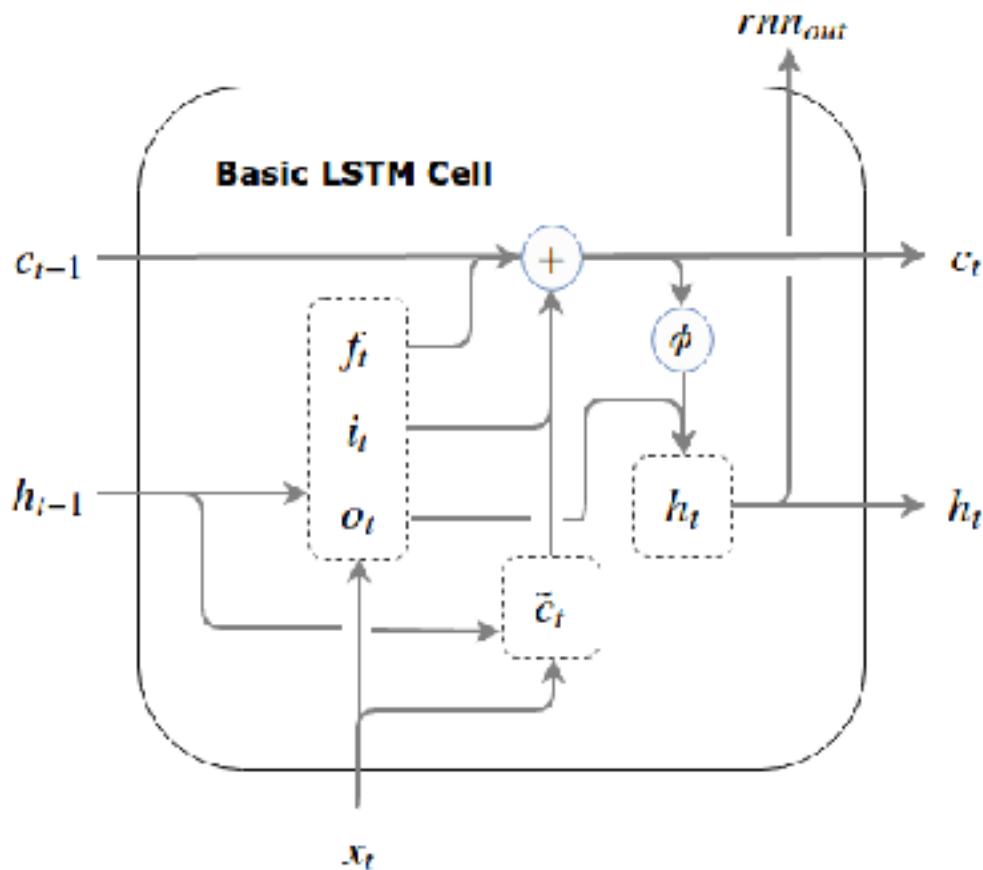
$$h_t = o_t \odot \phi(c_t)$$

get output from current state

$$rnn_{out} = h_t$$

LSTM Dropout

- LSTM in TensorFlow



$$\begin{aligned} i_t &= \sigma(W_i h_{t-1} + U_i x_t + b_i) \\ o_t &= \sigma(W_o h_{t-1} + U_o x_t + b_o) \\ f_t &= \sigma(W_f h_{t-1} + U_f x_t + b_f) \end{aligned}$$

Recurrent Dropout Input Dropout

The days of training **without** using **dropout** are **over**.

What to choose?

- There is no hard and fast rule
 - try both
 - Basic LSTM has had great success
 - GRU also sometimes is easier to train
 - you will see many variations
 - peephole LSTM
 - hierarchical LSTM
 - and many more...

Next time

- Recurrent Networks (two lecture agenda)
 - *Overview*
 - *Problem Types*
 - *Embeddings*
 - *Types of RNNs*
 - **Demo A**
 - **CNNs and RNNs**
 - **Demo B**
 - **course retrospective**

Lecture Notes for Machine Learning in Python

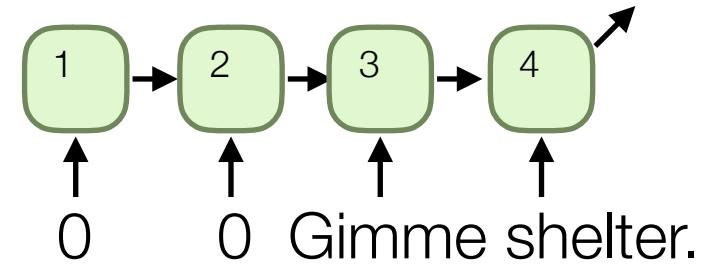
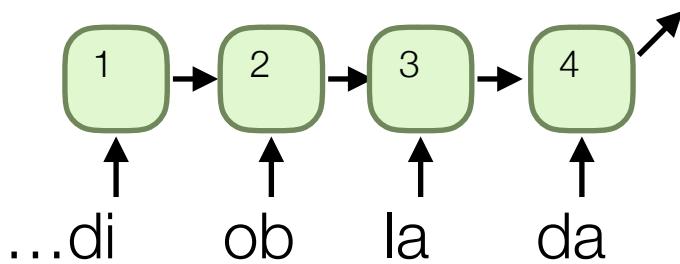
Professor Eric Larson
Final Lecture: RNN Demo

Lecture Agenda

- Logistics
 - RNNs due **Friday, Dec 1**
- Recurrent Networks (two lecture agenda)
 - *Overview*
 - *Problem Types*
 - *Embeddings*
 - *Types of RNNs*
 - Demo A
 - CNNs and RNNs
 - Demo B
 - course retrospective

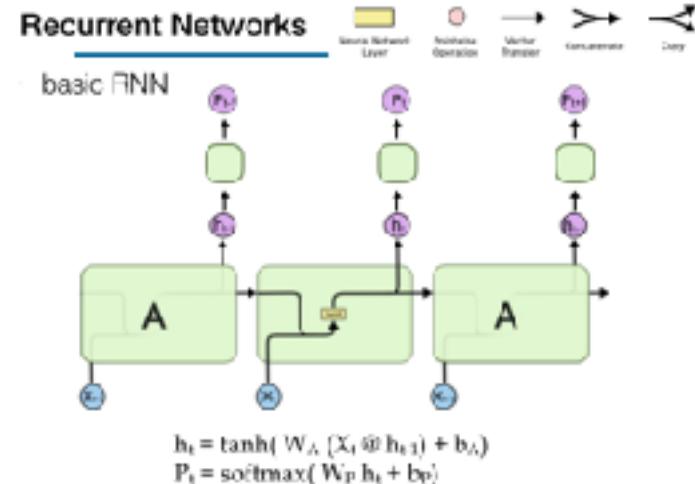
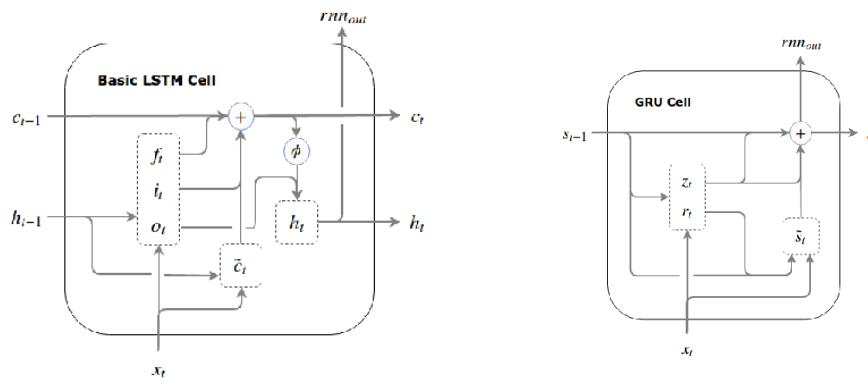
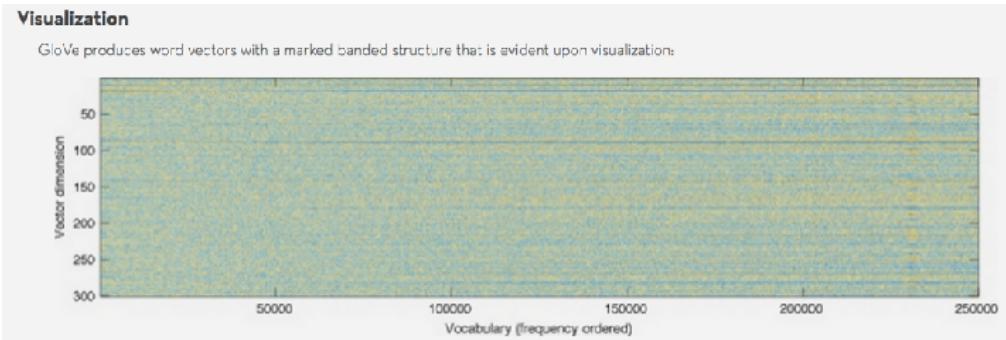
Last Time

- padding/clipping



Visualization

GloVe produces word vectors with a marked banded structure that is evident upon visualization:



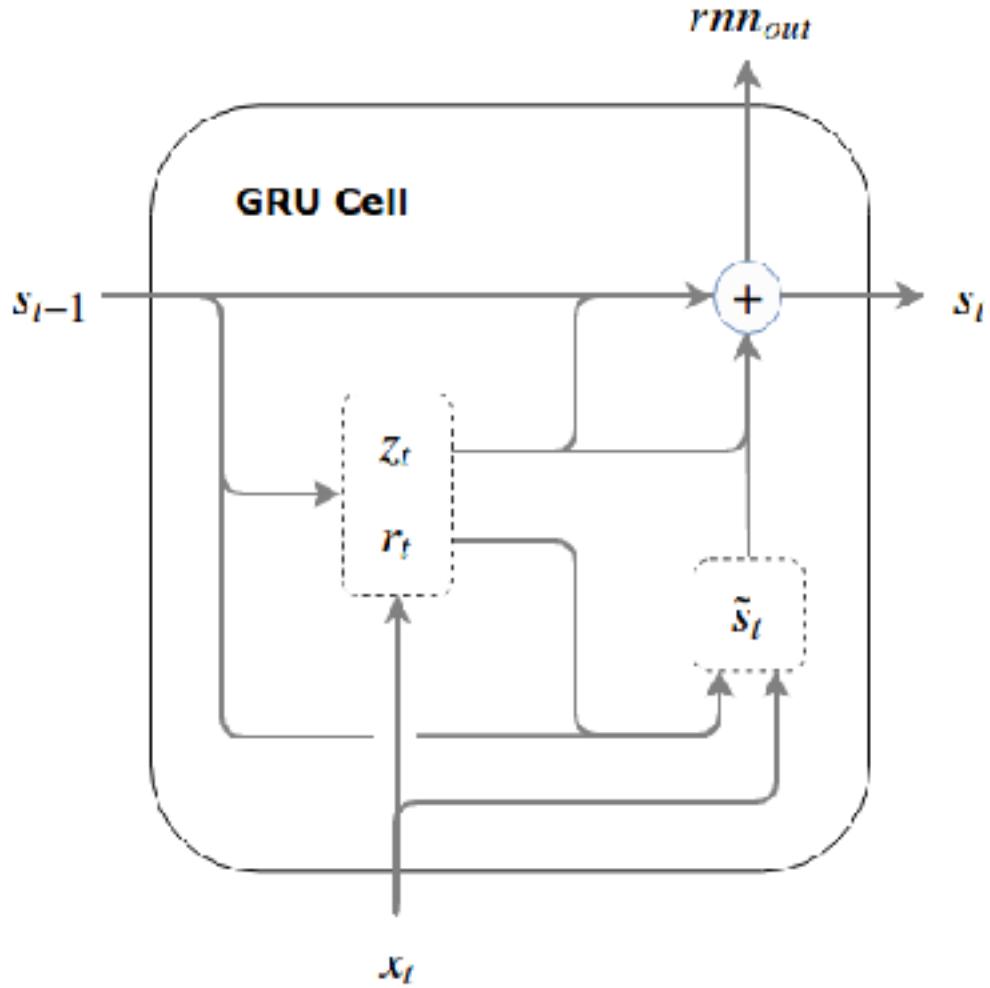
$$h_t = \tanh(W_h [x_t \otimes h_{t-1}] + b_h)$$

$$P_t = \text{softmax}(W_p h_t + b_p)$$

<http://colah.github.io/posts/2015-03-Understanding-LSTMs/>

Recurrent Networks: GRUs

- gated recurrent units



Selectivity controls, gates (**0 or 1**)

$$r_t = \sigma(W_r s_{t-1} + U_r x_t + b_r)$$
$$z_t = \sigma(W_z s_{t-1} + U_z x_t + b_z)$$

past state

current input

selectively remember

with influence

$$\tilde{s}_t = \phi(W(r_t \odot s_{t-1}) + U x_t + b)$$

$$s_t = z_t \odot s_{t-1} + (1 - z_t) \odot \tilde{s}_t$$

remember only past

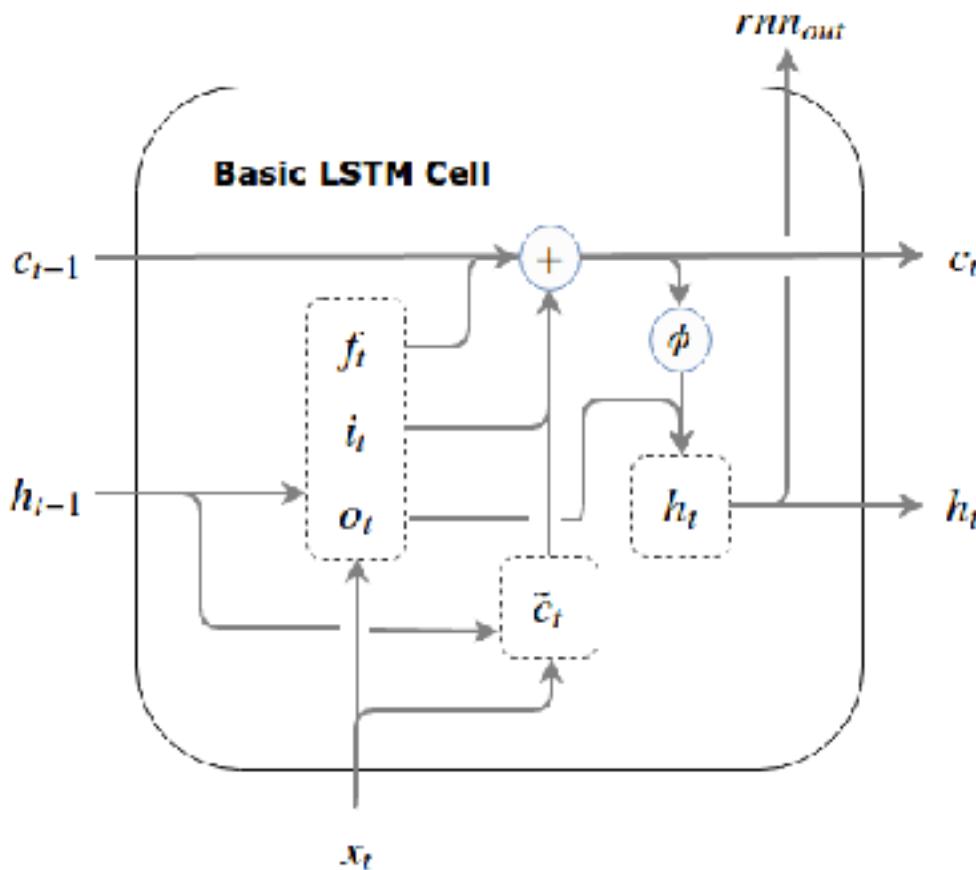
OR remember with input

σ = hard limit

\odot = elem. multiplication

Recurrent Networks: Gen 2 LSTM

- LSTM in TensorFlow



Selectivity controls

$$i_t = \sigma(W_i h_{t-1} + U_i x_t + b_i)$$

$$o_t = \sigma(W_o h_{t-1} + U_o x_t + b_o)$$

$$f_t = \sigma(W_f h_{t-1} + U_f x_t + b_f)$$

explicit remembering state

$$\tilde{c}_t = \phi(W h_{t-1} + U x_t + b)$$

$$c_t = f_t \odot c_{t-1} + i_t \odot \tilde{c}_t$$

remember
previous state

update with
new state

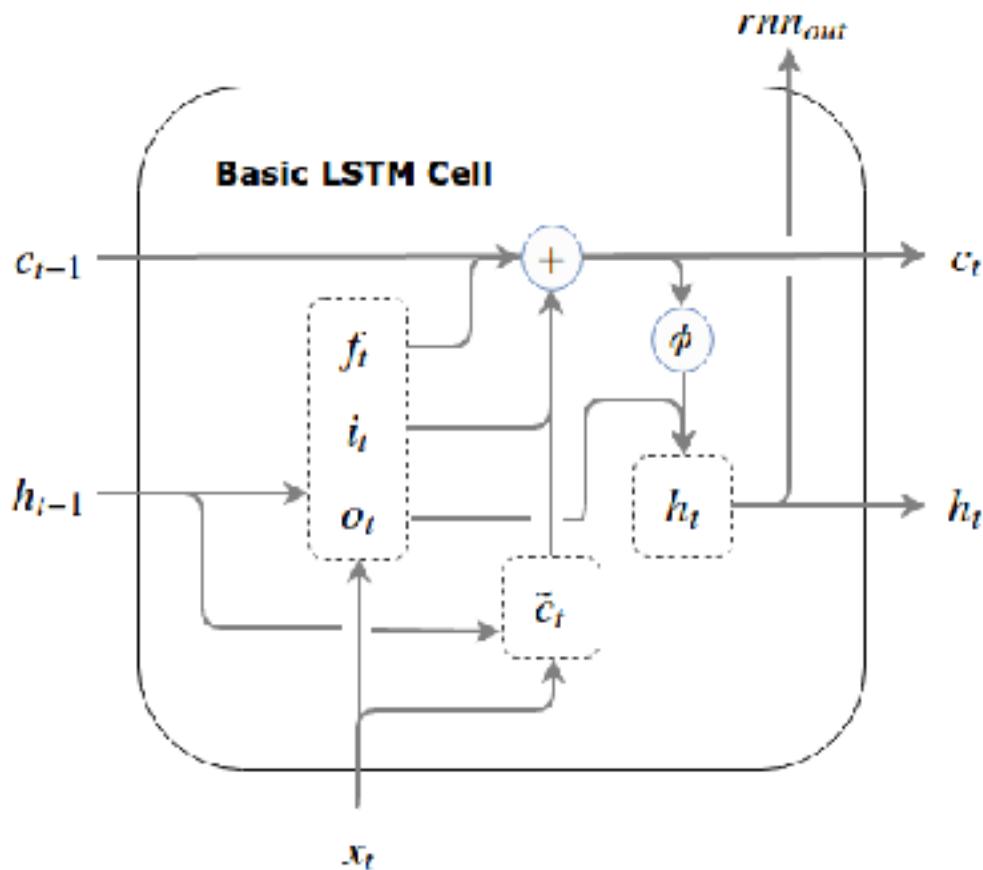
$$h_t = o_t \odot \phi(c_t)$$

get output from current state

$$rnn_{out} = h_t$$

LSTM Dropout

- LSTM in TensorFlow



$$\begin{aligned} i_t &= \sigma(W_i h_{t-1} + U_i x_t + b_i) \\ o_t &= \sigma(W_o h_{t-1} + U_o x_t + b_o) \\ f_t &= \sigma(W_f h_{t-1} + U_f x_t + b_f) \end{aligned}$$

Recurrent Dropout Input Dropout

The days of training **without** using **dropout** are **over**.

Recurrent Networks in Keras

Many to one:
Simple RNNs
GRUs
LSTMs



More examples:

<https://github.com/tensorflow/tensorflow/tree/r0.11/tensorflow/examples/skflow>

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>

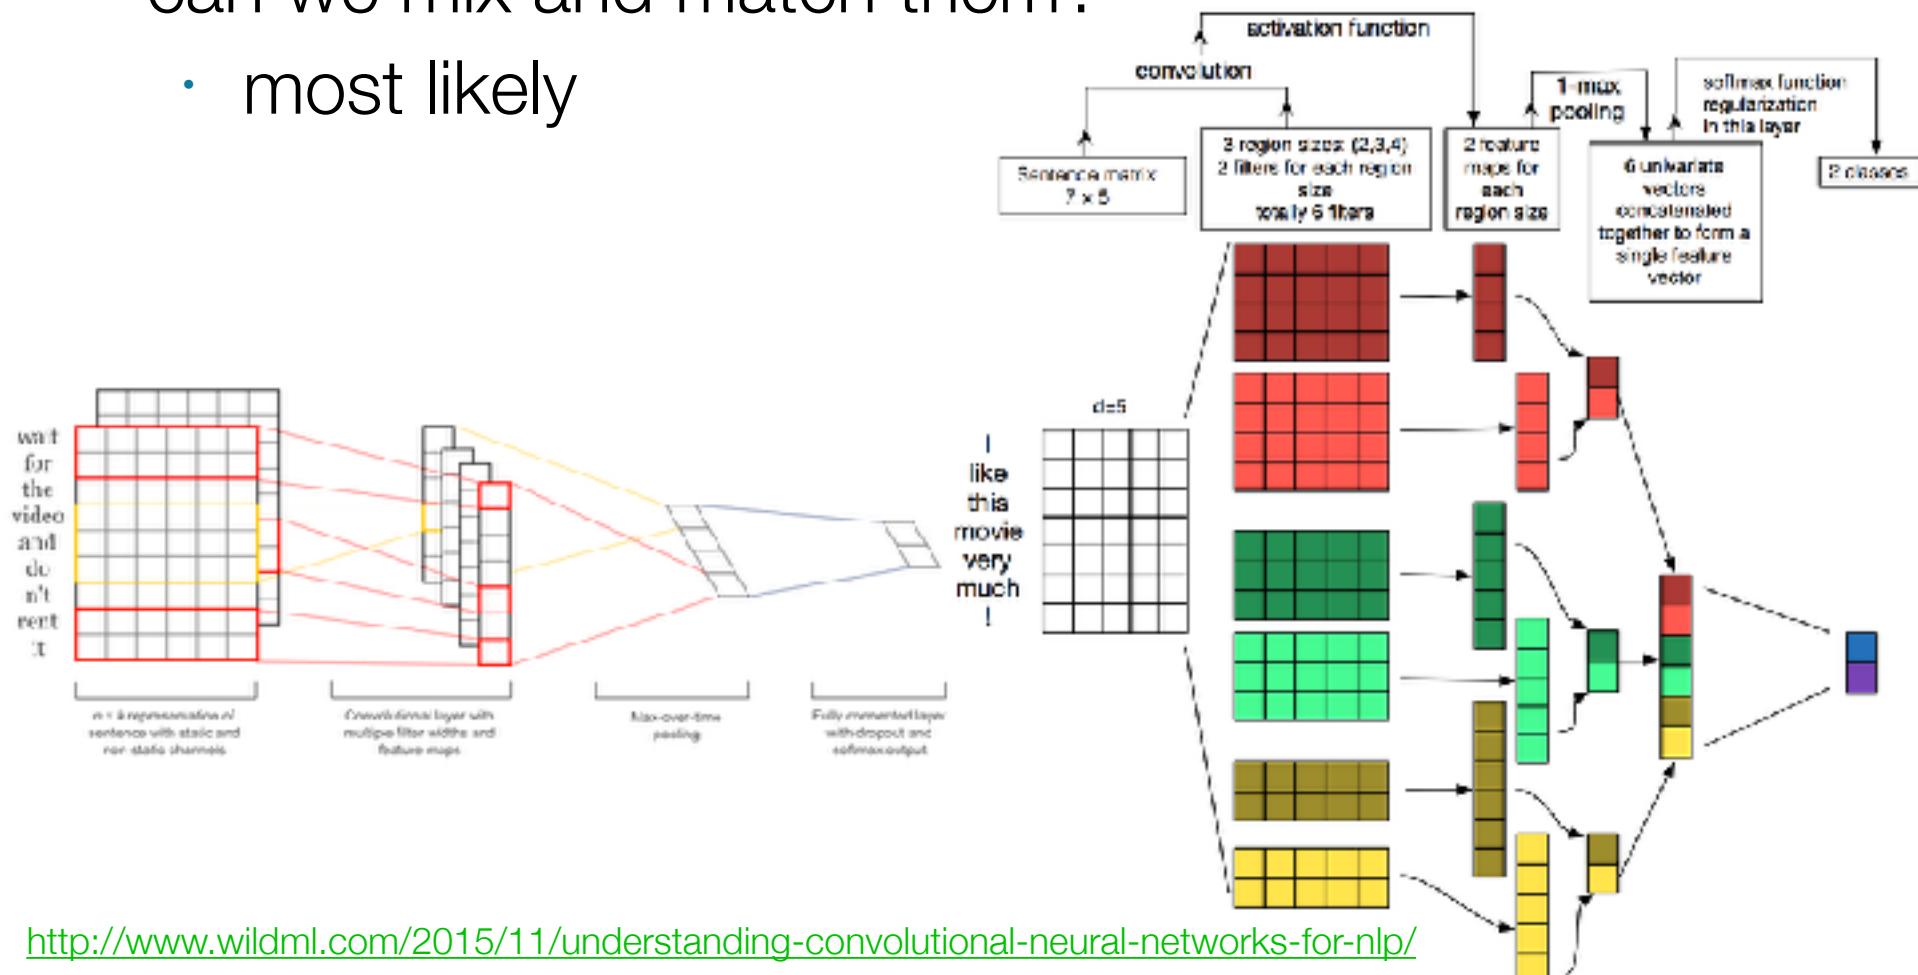
<http://machinelearningmastery.com/sequence-classification-lstm-recurrent-neural-networks-python-keras/>

Seq2Seq:

https://github.com/tensorflow/tensorflow/blob/r0.11/tensorflow/examples/skflow/neural_translator_word.py

CNNs and RNNs

- is an RNN similar to a CNN?
 - given fixed length sequences, probably
- can we mix and match them?
 - most likely



Recurrent Networks in Keras

Back to the CNN



More examples:

<http://www.wildml.com/2015/11/understanding-convolutional-neural-networks-for-nlp/>

Seq2Seq:

https://github.com/tensorflow/tensorflow/blob/r0.11/tensorflow/examples/skflow/neural_translation_word.py

Course Retrospective

Leading ML researchers issue statement of support for JMLR

- AI winters exist over-hyped p learning and the (and history)
- Formal methods in learning are needed
- At the end of the day, using back propagation
- Open source can still have a future without formal methods
 - <http://www.jmlr.org/mljresignation.html>

*From: Michael Jordan [mailto:jordan@CS.Berkeley.EDU]
 Sents: Monday, October 08, 2001 5:33 PM
 Subject: Letter of resignation from Machine Learning journal*

Dear colleagues in machine learning,

The forty people whose names appear below have resigned from the Editorial Board of the Machine Learning Journal (MLJ). We would like to make our resignations public, to explain the rationale for our action, and to indicate some of the implications that we see for members of the machine learning community worldwide.

The machine learning community has come of age during a period of enormous change in the way that research publications are circulated. Fifteen years ago research papers did not circulate easily, and as with other research communities we were fortunate that a viable commercial publishing model was in place so that the fledgling MLJ could begin to circulate. The needs of the community, principally those of seeing our published papers circulate as widely and rapidly as possible, and the business model of commercial publishers were in harmony.

Times have changed. Articles now circulate easily via the Internet, but unfortunately MLJ publications are under restricted access. Universities and research centers can pay a yearly fee of \$1050 US to obtain unrestricted access to MLJ articles (and individuals can pay \$120 US). While these fees provide income for institutions and individuals who can afford them, we feel that they also have the effect of limiting contact between the current machine learning community and the potentially much larger community of researchers worldwide whose participation in our field should be the fruit of the modern Internet.

None of the revenue stream from the journal makes its way back to authors, and in this context authors should expect a particularly favorable return on their intellectual contribution---they should expect a service that maximizes the distribution of their work. We see little benefit accruing to our community from a mechanism that ensures revenue for a third party by restricting the communication channel between authors and readers.

Chris Atkeson
 Peter Bartlett
 Andrew Barto
 Jonathan Baxter
 Yoshua Bengio
 Kristin Bennett
 Chris Bishop
 Justin Boyan
 Carla Brodley
 Claire Cardie
 William Cohen
 Peter Dayan
 Tom Dietterich
 Jerome Friedman
 Nir Friedman
 Zoubin Ghahramani
 David Heckerman
 Geoffrey Hinton
 Yoav Hirsh
 Tommi Jaakkola
 Michael Jordan
 Leslie Kaelbling
 Daphne Koller
 John Lafferty
 Sridhar Mahadevan
 Marina Meila
 Andrew McCallum
 Tom Mitchell
 Stuart Russell
 Lawrence Saul
 Bernhard Schoelkopf
 John Shawe-Taylor
 Yoram Singer
 Satinder Singh
 Padhraic Smyth
 Richard Sutton
 Sebastian Thrun
 Manfred Warmuth
 Chris Williams
 Robert Williamson

Thank you for a great semester!

- but it could **have been better** somehow, right?
 - if you went into a job interview for a data scientist position and were asked about **X**, that you did not feel strong in, what would **X** be?
 - i.e., what didn't translate well?
 - i.e., what are you **still not comfortable** with?
 - how could you learn better? same datasets to compare with others?
 - what should **not be** cut or changed?
 - what **skills** should students have **before this class** that I don't yet require?

Backup slides

TensorFlow

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>

```
with tf.variable_scope('rnn_cell'):
    W = tf.get_variable('W', [num_classes + state_size, state_size])
    b = tf.get_variable('b', [state_size], initializer=tf.constant_initializer(0.0))

def rnn_cell(rnn_input, state):
    with tf.variable_scope('rnn_cell', reuse=True):
        W = tf.get_variable('W', [num_classes + state_size, state_size])
        b = tf.get_variable('b', [state_size], initializer=tf.constant_initializer(0.0))
    return tf.tanh(tf.matmul(tf.concat([rnn_input, state]), W) + b)

state = init_state
rnn_outputs = []
for rnn_input in rnn_inputs:
    state = rnn_cell(rnn_input, state)
    rnn_outputs.append(state)
final_state = rnn_outputs[-1]

#logits and predictions
with tf.variable_scope('softmax'):
    W = tf.get_variable('W', [state_size, num_classes])
    b = tf.get_variable('b', [num_classes], initializer=tf.constant_initializer(0.0))
logits = [tf.matmul(rnn_output, W) + b for rnn_output in rnn_outputs]
predictions = [tf.nn.softmax(logit) for logit in logits]

# Turn our y placeholder into a list labels
y_as_list = [tf.squeeze(i, squeeze_dims=[1]) for i in tf.split(1, num_steps, y)]

#losses and train_step
losses = [tf.nn.sparse_softmax_cross_entropy_with_logits(logit,label) for \
          logit, label in zip(logits, y_as_list)]
total_loss = tf.reduce_mean(losses)
train_step = tf.train.AdagradOptimizer(learning_rate).minimize(total_loss)
```

recurrent networks

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>

```
def train_network(num_epochs, num_steps, state_size=4, verbose=True):
    with tf.Session() as sess:
        sess.run(tf.initialize_all_variables())
        training_losses = []
        for idx, epoch in enumerate(gen_epochs(num_epochs, num_steps)):
            training_loss = 0
            training_state = np.zeros((batch_size, state_size))
            if verbose:
                print("\nEPOCH", idx)
            for step, (X, Y) in enumerate(epoch):
                tr_losses, training_loss_, training_state, _ = \
                    sess.run([losses,
                             total_loss,
                             final_state,
                             train_step],
                             feed_dict={x:X, y:Y, init_state:training_state})
                training_loss += training_loss_
                if step % 100 == 0 and step > 0:
                    if verbose:
                        print("Average loss at step", step,
                              "for last 250 steps:", training_loss/100)
                training_losses.append(training_loss/100)
                training_loss = 0

    return training_losses
```

TensorFlow

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>

```
def train_network(num_epochs, num_steps, state_size=4, verbose=True):
    with tf.Session() as sess:
        sess.run(tf.initialize_all_variables())
        for idx, epoch in enumerate(gen_epochs(num_epochs, num_steps)):
            training_state = np.zeros((batch_size, state_size))
            for X, Y in epoch:
                tr_losses, training_loss_, training_state, _ = \
                    sess.run([losses,
                             total_loss,
                             final_state,
                             train_step],
                            feed_dict={x:X, y:Y, init_state:training_state})
```

TensorFlow (simplified)

<http://r2rt.com/recurrent-neural-networks-in-tensorflow-i.html>

```
cell = tf.nn.rnn_cell.BasicRNNCell(state_size)
rnn_outputs, final_state = tf.nn.rnn(cell, rnn_inputs, initial_state=init_state)

loss_weights = [tf.ones([batch_size]) for i in range(num_steps)]
losses = tf.nn.seq2seq.sequence_loss_by_example(logits, y_as_list, loss_weights)

x = tf.placeholder(tf.int32, [batch_size, num_steps], name='input_placeholder')
y = tf.placeholder(tf.int32, [batch_size, num_steps], name='labels_placeholder')
init_state = tf.zeros([batch_size, state_size])

x_one_hot = tf.one_hot(x, num_classes)
rnn_inputs = tf.unpack(x_one_hot, axis=1)

cell = tf.nn.rnn_cell.BasicRNNCell(state_size)
rnn_outputs, final_state = tf.nn.rnn(cell, rnn_inputs, initial_state=init_state)

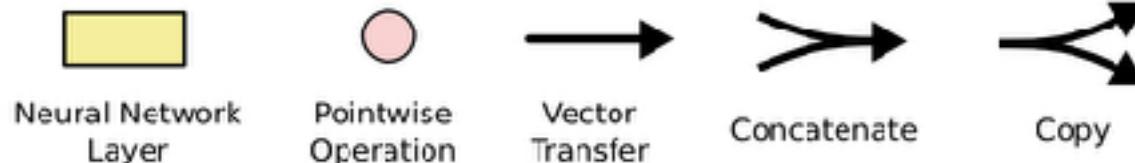
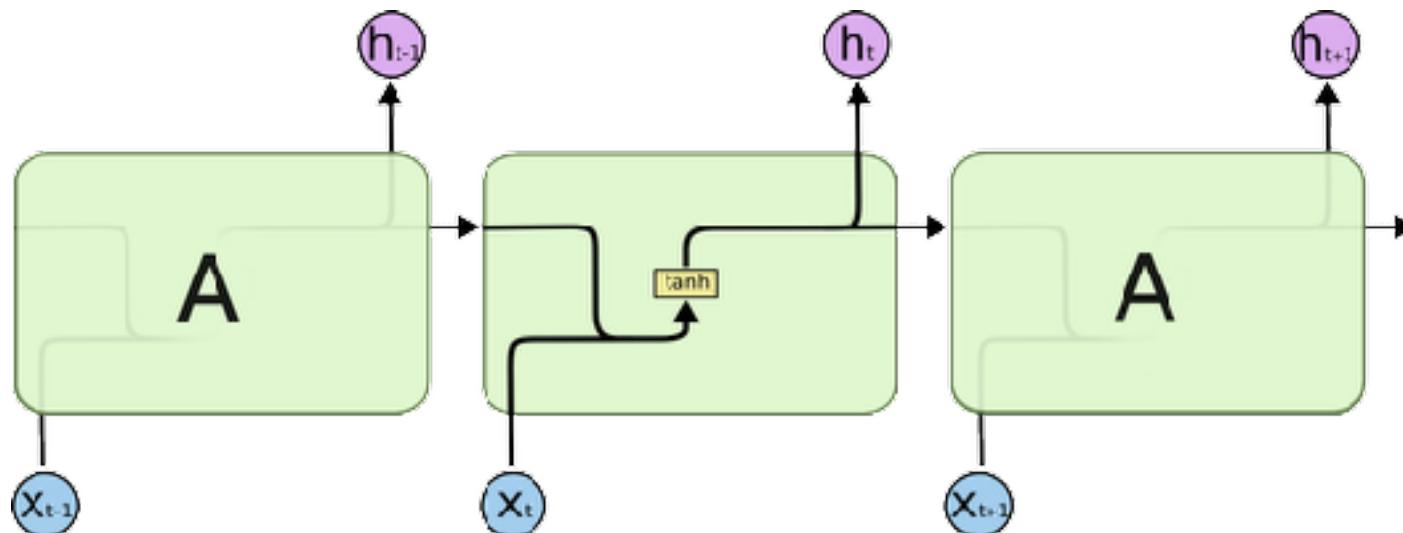
with tf.variable_scope('softmax'):
    W = tf.get_variable('W', [state_size, num_classes])
    b = tf.get_variable('b', [num_classes], initializer=tf.constant_initializer(0.0))
logits = [tf.matmul(rnn_output, W) + b for rnn_output in rnn_outputs]
predictions = [tf.nn.softmax(logit) for logit in logits]

y_as_list = [tf.squeeze(i, squeeze_dims=[1]) for i in tf.split(1, num_steps, y)]

loss_weights = [tf.ones([batch_size]) for i in range(num_steps)]
losses = tf.nn.seq2seq.sequence_loss_by_example(logits, y_as_list, loss_weights)
total_loss = tf.reduce_mean(losses)
train_step = tf.train.AdagradOptimizer(learning_rate).minimize(total_loss)
```

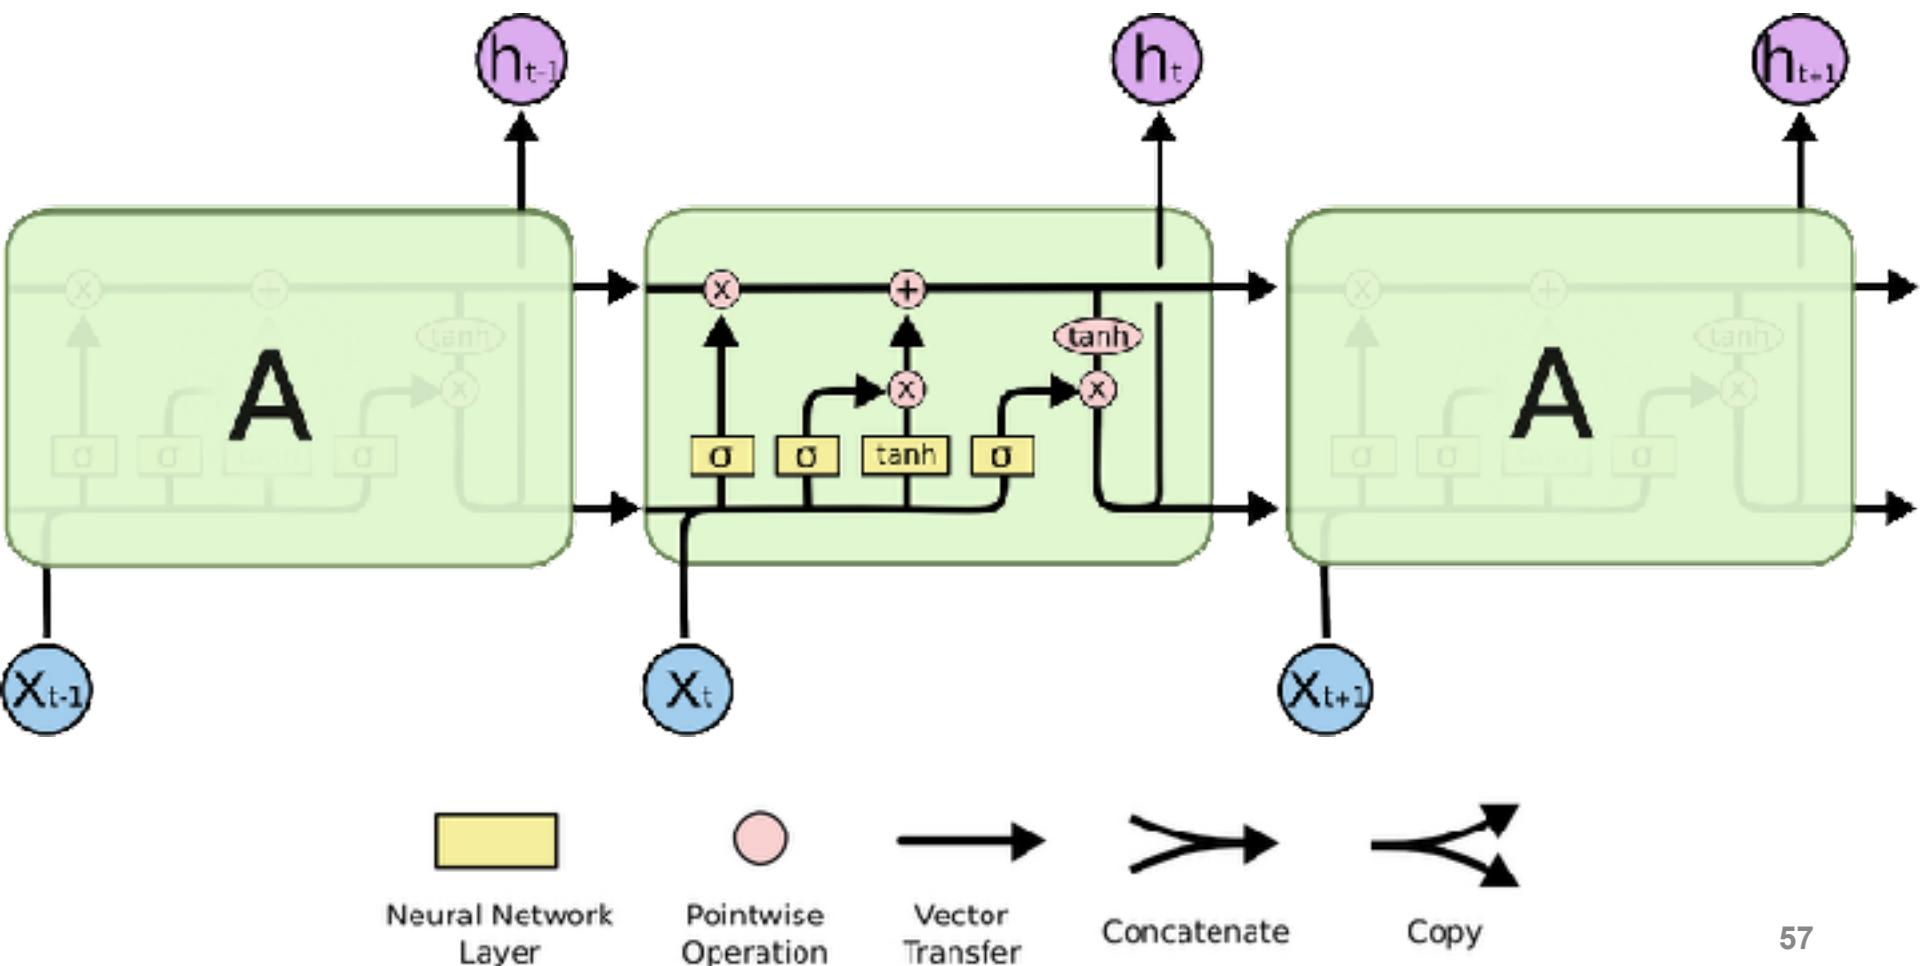
More Advanced Architectures

- **LSTM key idea:** limit how past data can affect output



More Advanced Architectures

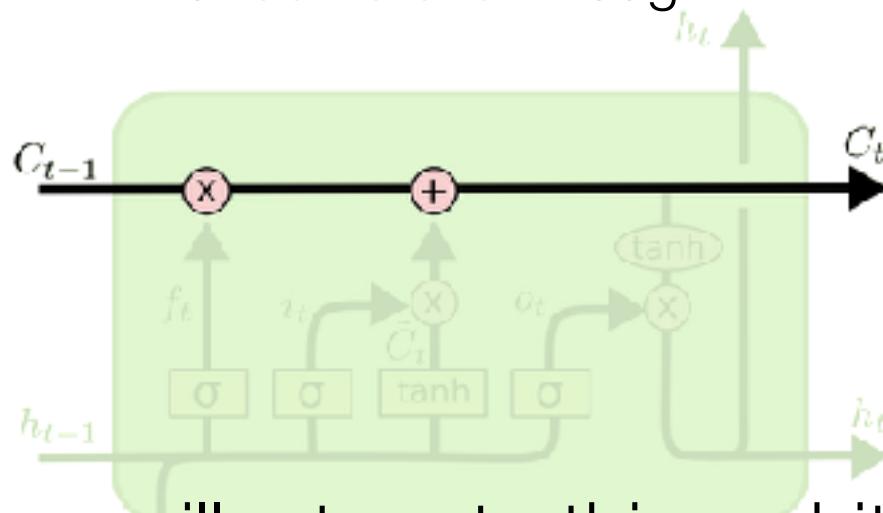
- **LSTM key idea:** limit how past data can affect output



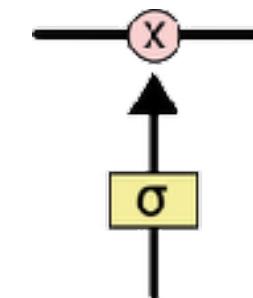
More Advanced Architectures

- **LSTM key idea:** limit how past data can affect output

let **cell state** through



potentially **forget past** inputs via “gate” σ



we will return to this architecture later, for now:

put it in long term memory 😂



Neural Network
Layer



Pointwise
Operation



Vector
Transfer



Concatenate



Copy